

Mind The Gap: Newsfeed Visualisation with Metro Maps

Damask Talary-Brown

Bachelor of Science in Computer Science with Honours
The University of Bath
2017

This dissertation may be made available for consultation within the University Library and may be photocopied or lent to other libraries for the purposes of consultation.

Signed:  Damas

Mind The Gap: Newsfeed Visualisation with Metro Maps

Submitted by: Damask Talary-Brown

COPYRIGHT

Attention is drawn to the fact that copyright of this dissertation rests with its author. The Intellectual Property Rights of the products produced as part of the project belong to the author unless otherwise specified below, in accordance with the University of Bath's policy on intellectual property (see <http://www.bath.ac.uk/ordinances/22.pdf>).

This copy of the dissertation has been supplied on condition that anyone who consults it is understood to recognise that its copyright rests with its author and that no quotation from the dissertation and no information derived from it may be published without the prior written consent of the author.

Declaration

This dissertation is submitted to the University of Bath in accordance with the requirements of the degree of Bachelor of Science in the Department of Computer Science. No portion of the work in this dissertation has been submitted in support of an application for any other degree or qualification of this or any other university or institution of learning. Except where specifically acknowledged, it is the work of the author.

Signed: 

Abstract

Contents

Abstract	i
List of Figures	vi
Introduction	1
1 Literature and Technology Review	2
1.1 Why don't we Understand the News?	2
1.1.1 Information Overload	3
1.1.2 Supporting Sensemaking	6
1.2 An Overview of Information Visualisation	7
1.2.1 Infovis for Sensemaking	8
1.2.2 Visual Metaphors	9
1.3 Metro Maps for Information Cartography	13
1.3.1 Coherence	14
1.3.2 Coverage	14
1.3.3 Connectivity	15
1.3.4 Limitations of Shahaf et al. [2012 <i>b</i> , 2013]	15
1.4 Towards Newsfeed Visualisation	17
1.4.1 Content Retrieval	17
1.4.2 Keyword Extraction	18
1.5 Evaluation Methods	19
1.6 Summary	20
2 Requirements	21
2.1 Project Scope	21
2.2 Requirements Gathering	22

2.3	Prioritisation	22
2.3.1	Categorisation	22
2.4	High-Level Design	23
2.5	Discussion	24
3	Design and Implementation	25
3.1	Introduction	25
3.2	Code Reuse	25
3.3	Article Retrieval	26
3.4	Keyword Extraction	26
3.4.1	Named Entity Recognition	27
3.4.2	Choosing a Term-Weighting Metric for Keyword Ranking	31
3.5	Graph Building	33
3.5.1	Extending the Definition of Coverage to Paths	33
3.5.2	Penalising Affinity	34
3.5.3	Coherence	35
3.5.4	Serialisation	35
3.6	Map Drawing	36
3.6.1	Stott's Criteria Metro Map layouts	37
3.6.2	Initial Force-Directed Positioning	38
3.6.3	Heuristic Refinement	41
3.7	Summary	43
4	Results and Discussion	44
4.1	UK Politics (BBC News)	44
4.2	Business (The Telegraph)	46
4.3	Sport (The Guardian)	48
4.4	Summary	50
5	Empirical Evaluation	51
6	Conclusions	52
A	Full Requirements Specification	59

- A.1 Functional Requirements 59
 - 1 Article Retrieval 59
 - 2 Keyword Extraction 59
 - 3 Graph Building 60
 - 4 Map Drawing 61
 - 5 Storage and Persistence 61
- A.2 Nonfunctional Requirements 62
 - 1 Security 62
 - 2 Software Quality 62
- B Implementation Raw Data 63**
 - B.1 Token/Entity Data 63

List of Figures

1.1	Dimensions of information overload, as defined by Ho and Tang [2001].	4
1.2	Similar visualisations from ThemeRiver [Havre et al., 2002] (left) and TIARA [Liu, Zhou, Pan, Qian, Cai and Lian, 2009] (right).	10
1.3	1930 London Underground map, designed by F. Stingemore [British Broadcasting Corporation, 2013].	12
1.4	1933 London Underground Map, designed by H. Beck [British Broadcasting Corporation, 2013].	12
1.5	A metro map [Shahaf et al., 2012b] covering the Greek Debt Crisis.	13
1.6	An incoherent chain with jittery transitions between topics (Chain A, left) alongside a more coherent chain of articles (Chain B, right). [Shahaf et al., 2012b]	14
1.7	Frequency spectra for ‘Quixote’ and ‘but’ in the first 50,000 words of <i>Don Quixote</i> . [Carpena et al., 2009]	19
2.1	A conceptual model of data flow between components of the system.	23
3.1	Stepping through the tokenisation process	26
3.2	Disambiguation threshold against pairs found and false positives (%)	30
3.3	Entities as a percentage of Tokens across 40 BBC Politics articles.	32
3.4	An example of three metro lines, each with high affinity.	34
3.5	‘Run-on’ affinity (left) is penalised more heavily than ‘one-off’ affinity (right).	34
3.6	Left-to-Right: A station (a), an interchange (b), a metro line (c), a link (d), a terminus (e) and a non-terminus station (f).	36
3.7	Octilinearity Calculation	39
3.8	Line Straightness Calculation	40
3.9	An example demonstrating the <i>snap-to-grid</i> process	42
3.10	An octilinear move which compromises the edge-length criterion	42
3.11	A non-octilinear move of an unplaced station	43

4.1	The front page BBC Politics article; drawn as (a) in Figure 4.2.	44
4.2	A metro map covering the BBC UK Politics RSS feed (March 16 th 2017) . . .	45
4.3	A metro map covering The Telegraph's Business RSS feed (March 17 th 2017) .	46
4.4	The redistribution of unbalanced edges along a line	47
4.5	A poor embedding of the map in Figure 4.3	47
4.6	A metro map covering the Guardian Sport RSS feed (March 16 th 2017) . . .	48
4.7	A metro map covering the Guardian Football RSS feed (March 17 th 2017) . . .	49

List of Algorithms

1	Finding unambiguous alias-entity pairs	28
2	Entity disambiguation with Knowledge Graph	30
3	Calculating Map Octilinearity	40
4	Calculating Map Line Straightness	41

Introduction

“The Press, Watson, is a most valuable institution, if you only know how to use it.”
— Sherlock Holmes, *The Adventure of the Six Napoleons*

Sir Arthur Conan Doyle

Aims

The aim of this project is the development of a tool which generates interactive metro maps of news from RSS feeds, with individual articles transformed into stations and common themes transformed into *metro lines*. My goal is to reduce the information overload experienced by news consumers, through the provision of contextual links and topic background within the visualisation of the feeds.

The resultant system is a news feed aggregator with graphically structured output, and to the best of my knowledge is first of its kind.

Outline

The structure of this dissertation is as follows:

- Chapter 1** provides an overview of the background literature upon which this project relies, including a detailed exploration of alternative visualisation formats.
- Chapter 2** describes the scoping of the system, the informal requirements gathering process undertaken, and the rationale behind the significant design decisions.
- Chapter 3** describes the algorithms and techniques chosen to transform the data from feed to visualisation and discusses how they were implemented in the system.
- Chapter 4** provides a set of example results generated by the system and discusses the effect of altering various parameters on the data.
- Chapter 5** describes the process by which the system was evaluated.
- Chapter 6** summarises both the contributions and limitations of this work, and provides a discussion on future research directions.

Chapter 1

Literature and Technology Review

1.1 Why don't we Understand the News?

The day the result of the 2016 United Kingdom EU Membership Referendum was announced, the @GoogleTrends Twitter account reported a 250% increase in searches for “What happens if we leave the EU?” Much like the case of David Leonhardt’s 2008 article in the New York Times in which began, “Raise your hand if you don’t quite understand this whole financial crisis,” national news commentary had focused on little else in the preceding months.

Some months after Leonhardt’s article was published, Journalism Professor Jay Rosen voiced his agreement with its premise in a blog post on the failure of journalism during the financial crisis; “there are certain very important stories – and the mortgage crisis is a good example – where until I grasp the whole I am unable to make sense of any part.” [Rosen, 2008]

It has become apparent that prolific coverage alone is not enough to engage and support the public in understanding the complexities of current events. Historically, news media has been limited in the volume of content it can produce by physical constraints such as printing costs, but the rise of the internet as a platform to deliver it has lead to an explosion of content, both through existing media channels and through competing social media websites and blogs.

The term *ambient news* was coined by Hargreaves et al. [2002] to describe the ubiquity of news in the current information landscape. Others have commented in a more critical light; describing the proliferation of competing news media as “as pervasive–and in some ways as invasive–as advertising.” [Nordenson, 2008, p.2]

In 2008, The Associated Press conducted an extensive field study into the news consumption habits of young adults. Among their key findings were three points which acutely summarise the news overload problem;

- **“Consumers are experiencing news fatigue.”**

The study found participants were debilitated and overwhelmed, and that their levels of dissatisfaction lead to a decrease in the effort they put into news acquisition. This is consistent with multiple other studies [Holton and Chyi, 2012, Purcell et al., 2010, Fischer and Stevens, 1991] which found participants across every demographic were overwhelmed by the amount of news content available to them and agreed that it prevented them exploring news on less familiar topics.

- **“Story resolution is key.”**

Participants’ consistent enjoyment of sports and entertainment news was due in part to the formulaic storytelling which characterises these types of journalism, with clear chronology to provide contextual back story. The feeling of enjoyment gained from reading procedural stories directly contrasts with what the same participants experienced reading World news, where they struggled to find resolution to stories which were unfolding at the time.

- **“Consumers want depth but aren’t getting it”**

It was observed that participants, in their efforts to discover *below-the-fold* content (defined in the context of the AP’s model, 2008, p.37) from particular headlines, often found themselves reading the same summary-level content from different news sources. It was recommended that news providers support this by “designing innovative formats and creating easier pathways to deep content.” [Associated Press and Context-Based Research Group, 2008, p.49]

Initially, the third point seems to be a direct contradiction to the first; we are overwhelmed by the volume of news we are exposed to, but we also crave more detail from the news we do consume. However, it brings to light the issue of information *quality* as a requirement of news consumers.

Journalism, and therefore its quality, can be viewed along a spectrum between two models; a model for the communication of facts, and a model for entertainment and storytelling. From the three points above, it is apparent that quality at both ends of the spectrum is being sought, since the desire for quality below-the-fold content is covered by the first model, and the desire for quality story resolution by the second.

1.1.1 Information Overload

News fatigue is a domain-specific type of information overload, a phenomenon formally defined as “when the information processing demands on time to perform interactions and internal calculations exceed the supply or capacity of time available for such processing” [Schick et al., 1990, p.206]. Information overload is a multifaceted problem which can be modelled as a combination of three contributing factors (Figure 1.1).

These factors correspond directly to the three points previously identified from the Associated Press and Context-Based Research Group study. High information quantity leads to news fatigue, information format determines the level of possible story resolution, and information quality determines how much depth a reader can gain from the news they consume. The authors did not find a single solution which could address all three factors, but they did identify information quantity as the most significant contributor to overload.

Bergamaschi et al. [2010] further decompose information quantity into spatial and temporal dimensions in the specific context of news articles. Spatial quantity refers to articles which are near-identical in terms of facts presented being published by different media outlets, and temporal quantity refers to articles on a single topic being published in quick succession over a short period of time.




Figure 1.1: Dimensions of information overload, as defined by Ho and Tang [2001].

Intuitively, the terms spatial and temporal quality seem illogically named, as a set of articles with high spatial quantity would cover a smaller area of information space and vice versa. High spatial quantity will therefore be referred to as *redundancy*, and high temporal quantity as *fragmentation*, since a sudden burst of articles published on the same topic suggests a currently unfolding story being told in parts.

To adequately determine the contributory factors relating to news fatigue, all four dimensions of information overload should be considered, and will therefore be explored in more detail in the following sections.

Information Quality

In the context of factual data rather than news specifically, Strong et al. [1997] defined information quality in terms of four components; intrinsic quality, accessibility quality, contextual quality and representational quality. If news can be rationalised to its core function as an interpretation of facts and other raw data such as images, then this same framework can be experimentally applied to news journalism in order to determine which factors could influence its quality.

Intrinsic quality is a measure of the accuracy, objectivity, believability and reputation of data. In the context of news, the first three factors would typically be true for all major news sources, and the reputation would be dependent on whether the article originated from a trusted source or not. Accessibility quality is less relevant to online news media, as it is concerned with data access and security. Contextual quality is the most relevant category in respect to news, concerning timeliness, amount of data, and value-added. In the news domain, this would mean an article's quality is dependent on its performance against a background of other articles; whether or not it contributes anything recent or previously unknown. Finally, representational quality is concerned with ease of understanding and interpretability, which are easily translatable concepts.

The implications of applying the Information Quality Framework [Strong et al., 1997] to news articles are that quality may be influenced by the reputation of the source, the timeliness of publication, value-added by the article (i.e. content which couldn't be derived from other sources) and the ease of understanding of the content.

Information Format

The domain of news articles is a more specific information space than that of documents in general, and by nature most news articles share some common formatting and structural elements such as headlines, timestamps, and relevant images. As a result of this, it is unlikely that any two articles from popular news providers would be diverse enough in content format to overshadow the information quantity problem.

Fragmentation (Temporal Quantity)

The rise of social media sites such as Twitter delivering news to consumers has lead to a high degree of news fragmentation, due to the constraints of the microblogging service's 140-character limit. 24-hour television news paved the way for new formats of real-time content delivery, and the ever-expanding network of online social media channels stepped up to deliver.

It logically follows from the fragmented nature of real-time news journalism that temporal quality suffers; stories are published and updated intermittently over short periods of time, meaning there is more content for the consumer to piece together in order to understand a story. The fragmentation is somewhat mitigated by Twitter's use of hashtags to denote a Tweet's topics. Hashtags help readers form a coherent and picture of unfolding events from the incremental contributions of thousands of participating users [Bruns et al., 2012].

Phuvipadawat and Murata [2010] developed a methodology to collect and group Tweets on breaking news topics, using hashtags for topic identification or *story-finding*, and grouping similar messages together to form a single news story. Their algorithm for similarity is a function of the tf-idf [Salton and Buckley, 1988] of the two messages and the number of named entities they have in common.

Redundancy (Spatial Quantity)

It is in the nature of news that newsworthy stories get repeated across multiple sources. When consumers read news on a particular from more than once source, it is likely that they will read variations on the same facts in multiple articles.

Attempts such as [Barzilay et al., 1999] have been made to synthesise summaries of collections of similar online documents, a practice here termed *information fusion*, with news articles from different sources being given as a specific use-case. However, the process of extracting common sentences between documents was in order to reformulate them into a single summary, rather than to determine the level of similarity between the documents.

A more relevant approach was presented by Pera and Ng [2008], who used the *title* and *description* attributes of elements in RSS feeds as content descriptors to mitigate the overhead

of processing entire documents for phrases. The content descriptors are then used to compute phrase n -grams as a measure of similarity between any two documents. The similarities in this case were used to remove subsumed articles and cluster non-redundant similar ones, in order to streamline feed content for readers.

It should be noted that there is an overlap between the notion of spatial quantity and one of the four influencing factors in information quality; contextual quality. If a feed contains two articles which state the same number of identical facts, they therefore contribute to information overload on both the qualitative and quantitative fronts.

Viewing the dimensions of overload from a news domain perspective, it is clear that (consistent with the findings of Ho and Tang [2001]) information quantity is the most relevant contributing factor in respect to fatigue, along factors influencing contextual quality such as value-added and timeliness. Any proposed solutions to the news overload problem should therefore address these factors first.

1.1.2 Supporting Sensemaking

Sensemaking is the basis for forming contextual knowledge; the process by which we incorporate new information into our existing cognitive frameworks, and how we go from reading something to understanding it [Pentina and Tarafdar, 2014]. In broader terms, Weick et al. describe sensemaking as “[being] about the question: What does an event mean? In the context of everyday life, when people confront something unintelligible and ask, ‘What’s the story here?’” [Weick et al., 2005, p.85]

This definition relates directly to the news overload problem because one component of sensemaking is the contextual story resolution. The Associated Press and Context-Based Research Group [2008] study identified news consumers are craving. It has also been observed that often readers are not interested in specific articles on a subject, and only the thematic content of the topic they belong to [Husin et al., 2014]. How then, do readers make sense of a collection of articles surrounding a particular topic?

When presented with a large document collection, Russell et al. [2006] found all of their subjects began by clustering the contents into groups which formed a heuristic representation or mental model, used to provide an overview. However, current information infrastructure has been criticised for not supporting the cross-correlation between connected news articles [Rennison, 1994].

Writing for the Columbia Journalism review in 2008, Nordenson outlined a suggestion for the new roles of journalism in the information era; “By linking stories to one another and to background information and analysis, news organizations help news consumers find their way through a flood of information that without such mediation could be overwhelming and nearly meaningless” [Nordenson, 2008, p.10].

Similarly, Pentina and Tarafdar [2014] make the recommendation in the context of contemporary media consumption that news providers should adapt to an environment of news overload by adding facilities enabling readers to categorise, sort and search news collections. Additional findings of this study suggested that the contextual background provided by having more detailed coverage aids the sensemaking process, as it helps users form links between

new information and their existing frameworks, but this presents an interesting conflict with the goal of reducing information overload when considering large collections of documents.

It is apparent that many recommendations have been made from within the field of journalism that at the point of delivery, news content should incorporate contextual links between related articles. This is important both from a sensemaking perspective to emphasise connections, and from an information overload perspective to help users find meaning in an inundated news landscape.

The news overload problem can now be reformulated with scope and detail: How can we display a collection of related news articles in such a way that users are not overwhelmed by unstructured content and are free to explore the underlying contextual pathways?

A simple starting point comes from a familiar idiom; a picture is worth a thousand words.

1.2 An Overview of Information Visualisation

Of course, a picture is not always worth a thousand words, particularly when the picture is unstructured and complex in its own right. However, a recognised and effective technique for bridging the gap between a set of data and a user's mental model and subsequent comprehension of the data is information visualisation, or infovis. [Yi et al., 2008, Havre et al., 2002]. This section provides a brief overview of a formative infovis taxonomy and uses the taxonomy to categorise appropriate visual models for newsfeed visualisation.

In his seminal paper on information visualisation, Shneiderman proposed a taxonomy for visualisations comprising seven data type abstractions, and seven tasks which are components of the visual information seeking mantra; "Overview first, zoom and filter, then details-on-demand." [Shneiderman, 1996, p.1]

Shneiderman's type abstractions are as follows:

1-dimensional	Linear data, where each datum is a string of characters.
2-dimensional	Planar data, e.g. layout diagrams, or clustered document collections.
3-dimensional	Physical objects or models of real-world entities, e.g. computer aided designs or medical imaging data.
Multi-dimensional	Any data where items with n attributes can be represented in n -dimensional space, e.g. relational databases, or feature vectors for classification.
Temporal	Data following a timeline, which is a subset of 1-dimensional data but was deemed important enough to warrant its own category. E.g. Project management data, or multimedia content timelines.
Tree	Hierarchical data where each datum has exactly one parent and zero or many children, e.g. document or directory structures.
Network	Related data, where each datum can have an arbitrary number of links to other data.

Because of the non-spatial nature of textual data, any visualisation of such data must involve some form of content abstraction and translation into a physical space [Wise et al., 1995]. These translations can result in data of arbitrary dimensionality, so a text corpus could fall into the 1-dimensional or multi-dimensional categories. News articles as a specific subset of textual data have certain metadata associated with them including dates, meaning they also fit the temporal type abstraction. In addition to this, if contextual links are considered part of the structure of the data, articles can be modelled as a network of connected nodes.

This ambiguity is not a failure of the taxonomy; Shneiderman stresses that composite categories are equally valid. However, the implications of this are that the most appropriate visualisation for a news corpus may itself be a composite of visualisations for any of its type abstractions, leaving an unfeasible number of possibilities to consider.

To reduce the scope of suitable visualisations, we return to the original problem of information overload. This time however, the aim is to minimise the overload from interpreting the model, rather than the overload from interpreting the data. Complex visualisations which require a considerable amount of effort to understand in their own right should be avoided where possible when reducing overload is the goal.

1.2.1 Infovis for Sensemaking

In addition to the insight that visualisation may be able to provide, there is evidence that visual metaphors better support the learning process and are more easily remembered than isomorphic text representations alone [O’donnell et al., 2002, Yen et al., 2012].

Yi et al. [2008] identified four overlapping infovis processes which describe how insight can be gained after sensemaking; *Provide Overview*, *Adjust*, *Detect Pattern*, and *Match Mental Model*. These four processes can be roughly mapped to Shneiderman’s high-level tasks, which are as follows:

Overview	Gain a birds-eye view of the entire collection, with the option to change the scale of the view by zooming or using fisheye magnification techniques.
Zoom	Gain a more detailed view of a portion of data or single datum while preserving the original sense of context.
Filter	Nondestructively remove uninteresting data points or groups from the view.
Details-on-Demand	Gain additional insight into one or more data points by selecting particular elements.
Relate	View and explore relationships between elements.
History	If necessary, undo previous actions to return to the a view of the data.
Extract	Export selected data, preserving the format, for uses such as “sending by email, printing, graphing, or insertion into a statistical or presentation package” [Shneiderman, 1996, p.5].

The *Provide Overview* process allows a reader to recognise what they know and what they don't know from the information they are processing. The corresponding task in [Shneiderman, 1996] is *Overview*.

Adjust allows them to change the level of abstraction or field of selection of that information. This corresponds to *Zoom* and *Filter* in Shneiderman's task model.

The *Detect Pattern* procedure is where structure and trends are found (whether expected or otherwise). Coupled with *Match Mental Model*, where the links are formed between the new data and the users' existing cognitive frameworks, this corresponds to *Relate*.

At this point, Shneiderman's taxonomy diverges from the processes of [Yi et al., 2008], as Yi et al. are concerned with the cognition enabled by visualisation, whereas Shneiderman additionally considers other use cases for visualisations, such as querying and sharing.

From these two models, it is apparent that certain views and functions are crucial for tools which use visualisation to support the sensemaking process; an high-level overview visualisation which emphasises links between data, the ability to adjust scope to show more or less detail, and the ability to filter information of specific interest within the dataset.

1.2.2 Visual Metaphors

Eppler defines visual metaphors as “a graphic structure that uses the shape and elements of a familiar natural or man-made artefact or of an easily recognizable activity or story to organize content meaningfully and use the associations with the metaphor to convey additional meaning about the content.” [Eppler, 2006, p.203] This definition highlights the main advantage to using visual metaphors; users are intuitively familiar with how they present and structure information.

The use of preexisting visual metaphors—specifically those with which a large number of people will already be familiar—has been shown to support readers' comprehension, as it requires both significant time and effort for a reader to interpret visual metaphors which are new to them [Ziemkiewicz and Kosara, 2009].

In a previous paper, Eppler also describes six advantages of visual metaphors specifically for the transfer of knowledge: “(1) to motivate people; (2) to present new perspectives; (3) to increase remembrance; (4) to support the process of learning; (5) to focus the attention of the viewer and (6) to structure and coordinate communication.” [Burkhard, 2004, p.2, citing [Eppler, 2004]]. These are all desirable attributes, but motivation and support in the learning process are particularly relevant in addressing news fatigue.

Examples of visual metaphors commonly used to represent collections of data include calendars, bookshelves, timelines, maps and other schematics. Metaphors based on physical objects do not have to be visually skeuomorphic to be effective, but to avoid misinterpretation, there should be a match between the underlying structure of the metaphor and the underlying structure of the data. Two common and highly structured visual metaphors will be explored in the following sections in the context of potential for news representation; timelines and schematic maps.

Timeline Visualisations

Chronological ordering is an important characteristic of news articles and should be preserved in any visualisation of news data as it provides a natural ordering [Binh Tran, 2013]. Perhaps the simplest visual metaphor for a collection of dated documents is the timeline.

Nguyen et al. [2014b] explore the role of timelines in the sensemaking process, emphasising that the interactions supported by such visualisations should be as intuitive as possible in order to not disrupt users' trains of thought, and should be tightly coupled with other elements of the sensemaking process so temporal connections are not viewed solely in isolation.

Criticisms of previous timeline visualisations made by the authors are that linear layouts are often too simple for the data they represent, and that a lack of automatic layout generation results in additional manual work for the user.

The colouring technique used to distinguish sets of related events within a single timeline is a flexible extension to [Nguyen et al., 2014a], where the authors coloured events belonging to multiple sets with a gradient composed of the colours of both sets. However, the gradient approach presented in [Nguyen et al., 2014b] does not scale to events belonging to more than two sets, since the colour grouping restricts the number of possible intersections of each set. From a news storyline perspective, this would place an upper limit of two on the number of possible topics a story could belong to, which is a low constraint for all but the highest level topics.

Singh et al. [2015] designed a prototype for generating annotated timelines based on the Wikipedia entries long-running news stories. The use of Wikipedia rather than newsfeeds meant their document retrieval model was heavily dependent on Wikipedia’s structure, but it also afforded a huge wealth of contextual information that made such detailed annotations possible. Not all stories are long running however, so while this would be useful as a retrospective tool it would be impossible to generate timelines in the same way for news articles which did not already belong to a long-running chain of events.

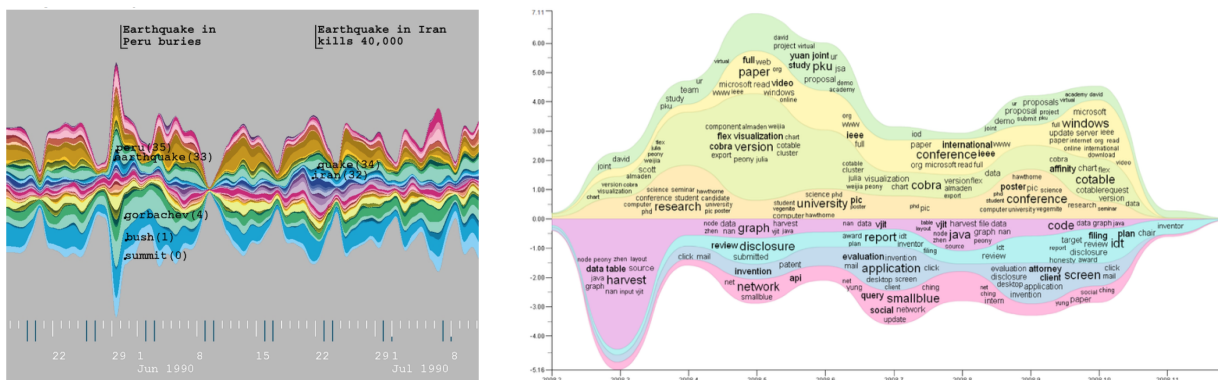


Figure 1.2: Similar visualisations from ThemeRiver [Havre et al., 2002] (left) and TIARA [Liu, Zhou, Pan, Qian, Cai and Lian, 2009] (right).

Both ESTHETE [Goyal et al., 2013] and nReader [Wang et al., 2006] present timeline-centric views for collections of news articles based on underlying graphs of relationships between the articles. However, in both cases, the graph structure was not part of the final visualisation, so

connections between entities were displayed in purely textual forms. In contrast, ThemeRiver [Havre et al., 2002] introduces a novel view on topic frequency along the time axis to show thematic change over time within a collection of documents, similar to a smoothed histogram. This view, while useful for large document collections which span weeks or months, would be less suited to displaying emerging news trends over shorter periods.

TIARA [Liu, Zhou, Pan, Qian, Cai and Lian, 2009] whose authors cite ThemeRiver as an influencing design (see figure 1.2), displays a similar shaped graphical output but performs more detailed textual analysis, and displays related keywords in the output. Both visualisations support simple zooming and panning, but suffer from the same limitations on visualising topic connectivity as [Nguyen et al., 2014a].

Taking into consideration both the critiques of oversimplification identified in [Nguyen et al., 2014b] and the physical limitations highlighted in [Nguyen et al., 2014a, Havre et al., 2002], it is clear that timelines may not be the most appropriate visual metaphor for news visualisation, especially not for highly connected events and topics. However, for more linear storylines which span fewer categories or topics, a visualisation such as [Nguyen et al., 2014b] could be used, for example as part of Shneiderman’s *Zoom and Filter* task where the dataset is pruned.

Topological and Schematic Map Visualisations

The use of cartographic representations for abstract objects and the relationships between them is such a common and natural metaphor that it is often unnoticed in digital contexts. Maps take advantage of humans’ natural ability to perceive and organise in a spatial context, and [the fact that we live in a spatial world] “leads naturally to metaphors that provide cues for orientation and navigation” [Old, 2002, p.2].

Topological maps¹ are a specific class of map which abstract away detail so that only significant features within desired subsets of the mapped dataset remain; from a news overload perspective, this is a highly desirable property. These maps are often used to visualise networks and are represented as schematics, where elements on the map are transformed into abstract visual representations for ease of understanding. Today, the most recognisable examples of topological maps are transit maps. Divorced from the strictness of geographical accuracy and scale, emphasis is instead placed on the usability of the map for planning and the understanding of relative positioning which the maps enable [Hochmair, 2009].

Figures 1.3 and 1.4 show the geographically accurate 1930 London Underground map designed by Frederick Stingemore, and the iconic 1933 redesign by Henry Beck, which was based on the concept of an electrical circuit diagram [Transport For London, 2014]. While Beck’s simplification of the structure of the map was seen as radical and even controversial at the time, the hallmarks of his design are now recognisable not just in other schematic maps, but in posters, infographics and diagrams across various domains.

In terms of usability, simple timeline visualisations can provide the chronological ordering which is intrinsic within collections of news articles. However, schematic maps may be able to represent complex relationships between topics in a structure which is linear but has an additional dimension, to prevent linearity becoming a design or usability constraint.

¹Not to be confused with *topographic* maps, which represent relief and other geographic features of physical regions at a large scale and in fine detail.

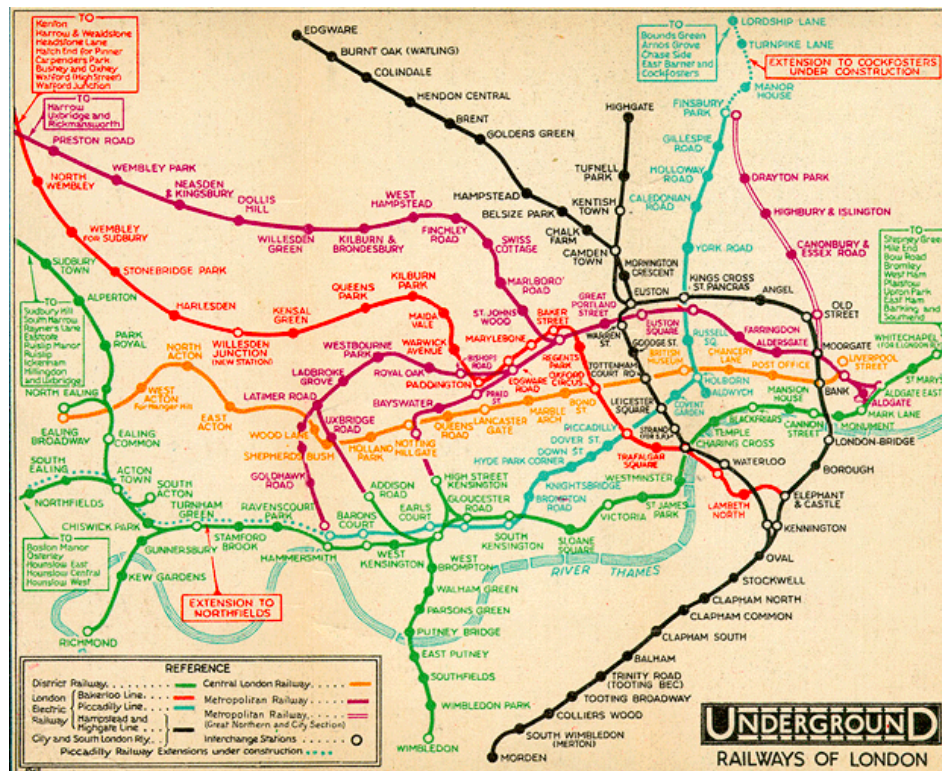


Figure 1.3: 1930 London Underground map, designed by F. Stingemore [British Broadcasting Corporation, 2013].

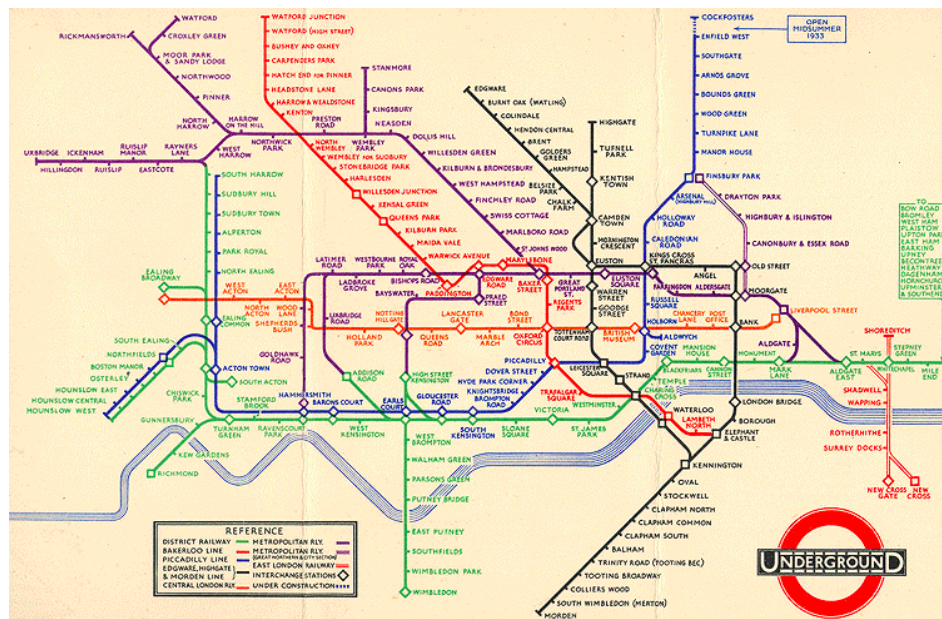


Figure 1.4: 1933 London Underground Map, designed by H. Beck [British Broadcasting Corporation, 2013].

1.3 Metro Maps for Information Cartography

Significant work in the area of information cartography has been undertaken by Shahaf et al. [Shahaf and Guestrin, 2010, Shahaf et al., 2012b,a, 2013], in the domains of both news and science through the visualisation of article and journal data on metro maps.

Shahaf et al. [2012b] chose the metaphor as a base for their visualisation to address the fact that previous timeline-based summarisation systems could only represent simple linear stories; “In contrast, complex stories display a very non-linear structure: stories split into branches, side stories, dead ends, and intertwining narratives.” [Shahaf et al., 2013, p.1]

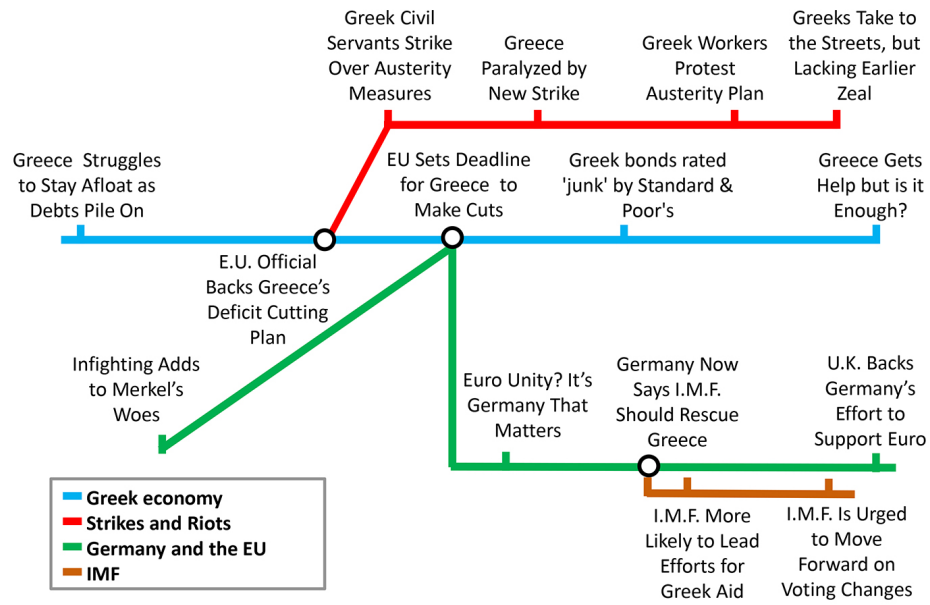


Figure 1.5: A metro map [Shahaf et al., 2012b] covering the Greek Debt Crisis.

Even in a scientific context, this was not the first time the abstract visualisation potential of the metro map had been noted; “The usefulness of the metro map as a metaphor is somewhat limited to simple examples by the time required to manually produce these maps. As such they are generally only useful for applications that do not change frequently. This limitation could be removed by quality methods for the automatic drawing of metro maps from abstract data.” [Stott, 2008, p.54]

In this section, the formalisation of the metro map metaphor, its associated characteristics, and its limitations will be discussed.


Definition 1. *Metro Map* [Shahaf et al., 2012b]: A metro map \mathcal{M} is a pair (G, Π) , where $G = (V, E)$ is a directed graph and Π is a set of paths, or metro lines in G . Each $e \in E$ must belong to at least one metro line.

A previously published method [Shahaf and Guestrin, 2010] for linking together chains of articles was discussed, and an objective function was created to formalise the characteristics of a ‘good’ metro map. The function defined was a composite based on three important

characteristics, all of which are broadly applicable to the visualisation of any similar corpora; coherence, coverage, and connectivity.

1.3.1 Coherence

Let \mathcal{D} be a set of articles, and \mathcal{W} be a set of words or phrases, such that each article is a subset of \mathcal{W} . A *coherent* chain of articles through \mathcal{D} is one where transitions between documents are smoothed by common overlapping keywords from \mathcal{W} , creating a better narrative flow [Shahaf and Guestrin, 2010] as depicted in Figure 1.6.



A bar corresponds to the presence of a word in the article above it.
The titles of the articles which made up the two chains were as follows:

Chain A (left)	Chain B (right)
Europe weighs possibility of debt default in Greece	Europe weighs possibility of debt default in Greece
Why Republicans don't fear a debt default	Europe commits to action on Greek debt
Italy; The Pope's leaning toward Republican ideas	Europe union moves towards a bailout of Greece
Italian-American groups protest 'Sopranos'	Greece set to release austerity plan
Greek workers protest austerity plan	Greek workers protest austerity plan

Figure 1.6: An incoherent chain with jittery transitions between topics (Chain A, left) alongside a more coherent chain of articles (Chain B, right). [Shahaf et al., 2012b]

Coherence, intuitively, seems to be closely linked to idea of story resolution detailed in Section 1.1.1. This presents a question which could later be explored further; does forming coherent chains of articles provide the story resolution that participants in the Associated Press study were so desperately seeking from current events journalism?

1.3.2 Coverage

As in the previous section, let \mathcal{D} be a set of articles, and \mathcal{W} be a set of words or phrases of which the articles are composed. The coverage function for a word in a given document $d_i \in \mathcal{D}$ specified in Equation 1.1 can be quantified using any measure of how well d_i covers w , for example $\text{tf-idf}(w, d_i, \mathcal{D})$ (See Equation 1.6) [Shahaf et al., 2012b].

$$\text{cover}_{d_i}(w) : \mathcal{W} \rightarrow [0, 1] \quad (1.1)$$

Extending the notion of coverage to maps—which can be abstracted to sets of documents—introduces the idea of *diversity*. If a map already contains documents which for a sufficient coverage for some word w , then there is nothing to be gained by adding another document to \mathcal{D} which has high coverage of w alone. This relates back to the principles of spatial and information quality discussed in Section 1.1.1, especially the importance of value-added by

every individual document in a collection. In this case, maps which cover a maximal number of $w \in \mathcal{W}$ should be preferential. A simple additive definition for map coverage such as Equation 1.2 [Shahaf et al., 2012b] would not reward this kind of diversity;

$$cover_{\mathcal{M}}(w) = \sum_{d_i \in docs(\mathcal{M})} cover_{d_i}(w) \quad (1.2)$$

Therefore, an alternative definition for map coverage was chosen, which will not increase significantly if another document which covers an already covered feature is added to \mathcal{D} (Equation 1.3 [Shahaf et al., 2012b]).

$$cover_{\mathcal{M}}(w) = 1 - \prod_{d_i \in docs(\mathcal{M})} (1 - cover_{d_i}(w)) \quad (1.3)$$

Finally, the definition of map coverage is extended to the coverage of the corpus \mathcal{D} , rather than just single features. If each feature is weighted according to frequency, then for each $w \in \mathcal{W}$ we have some λ_w . The coverage of a corpus \mathcal{D} by a metro map \mathcal{M} can then be defined as in Equation 1.4 [Shahaf et al., 2012b].

$$Cover(\mathcal{M}, \mathcal{D}) = \sum_{w \in \mathcal{W}} \lambda_w cover_{\mathcal{M}}(w) \quad (1.4)$$

1.3.3 Connectivity

The final property is the most simply defined; the connectivity of a metro map is the number of paths in Π which intersect [Shahaf et al., 2012b].

$$Connectivity(\mathcal{M}) = \sum_{i < j} \mathbb{1}(p_i \cap p_j \neq \emptyset) \quad (1.5)$$

1.3.4 Limitations of Shahaf et al. [2012b, 2013]

Corpus

Perhaps the biggest limitation of the system developed in [Shahaf et al., 2012b, 2013] is the nature of the corpus \mathcal{D} ; it is a fixed dataset, meaning users can only query it for certain past events with no way of specifying a different corpus themselves.

From a historical reference perspective the output generated based on certain queries is interesting when compared to the output an expert would select as important to the narrative, but it is not possible to use the system as a replacement to a newsfeed aggregator.

Graph Layout Aesthetic Principles

From a usability perspective, a second limitation of the work of Shahaf et al. is the lack of focus given to the desirable aesthetic properties of transit maps.

A formative empirical study [Purchase et al., 1997] on how graph layout affects usability identified a collection of five measurable aesthetic principles from previous research, which aid human understanding when reading graphs:

- Minimise the number of line bends [Tamassia, 1987];
- Minimise the number of edge crossings [Ferrari and Mezzalana, 1970];
- Preserve any underlying symmetry in the structure of the graph [Lipton et al., 1985];
- Draw orthogonally where possible [Tamassia, 1987];
- Maximise the minimum angle between incident edges for each node [Garg and Tamassia, 1994].

These criteria, when numerically calculated and weighted relative to one another, allow the aesthetic quality of any graph to be evaluated. This is particularly useful for graphs generated by automatic layout algorithms, which aren't as much a product of human design intuition as their manually designed counterparts.

The origins of the principles above predate much of the early infovis research because several of them [Tamassia, 1987, Ferrari and Mezzalana, 1970] were formalised as guidelines for the design of electronic circuits.

In a later study, Purchase evaluated information-finding task performance with a set of graphs which varied the above principles, to establish which the most and least significant aesthetics were for graph usability. The results of the study were that maximising incident edge angles and orthogonality did not lead to better performance, preserving symmetry and minimising line bends were somewhat important, and minimising edge crossings was the most significant influencing factor for performance [Purchase, 1997].

Metro maps, however—particularly those representing non-physical data such as in the previous section—do not share all the structural principles of the wider set of directed graphs. Consequentially, Stott et al. developed a set of criteria for metro maps specifically, including criteria for the labelling of stations. The criteria are as follows [Stott et al., 2011]:

- **Angular Resolution Criterion:** Maximise the angle between incident edges at each node. This criterion is also the fifth principle in [Purchase et al., 1997].
- **Edge Length Criterion:** All edges on the map should be approximately equal in length.
- **Balanced Edge Length Criterion:** The length of edges incident to a given node should be approximately equal.
- **Edge Crossings Criterion:** Crossings should be minimised. This criterion is also the second principle in [Purchase et al., 1997], also identified as the most important.
- **Line Straightness Criterion:** Edges on the same metro line should be collinear, i.e. they should form a 180° line through every station that the line passes through. This criterion relates closely to the first principle in [Purchase et al., 1997].
- **Octilinearity Criterion:** Edges should be drawn at multiples of 45° . This criterion is an extension to the fourth principle in [Purchase et al., 1997], as orthogonality is here analogous to rectilinearity.

Evaluating Figure 1.5 according to the list above, we observe that despite only containing 16 stations, the map unnecessarily violates five of Stott et al.’s six criteria—all except edge crossings.

Therefore, to improve the usability of the metro maps drawn in [Shahaf and Guestrin, 2010, Shahaf et al., 2012b,a, 2013], it would be advisable to attempt to satisfy more of the aesthetic principles for graphs and metro maps described in [Purchase et al., 1997, Stott et al., 2011].

1.4 Towards Newsfeed Visualisation

The fact that news articles form a fairly narrow class of document is an advantage from a visualisation design perspective, due to the common elements they share. Articles published by commercial news producers typically contain:

- A headline;
- A description, or *subhead*;
- A publish date;
- One or more categories to which the article belongs.

These attributes are useful for visualisation, since creating a spatial representation from text requires documents to be represented as vectors in high-dimensional feature space [Wise et al., 1995], and the presence of existing attributes makes articles more inherently comparable than their unstructured contents would be.

There is also a well-known existing standard for publishing links to articles with their metadata for use by other applications; RSS.

1.4.1 Content Retrieval

The de-facto web format for feed publishing is RSS (Rich Site Summary, or Really Simple Syndication.) The rise of the internet as a news platform has lead to many readers finding the most efficient method of reading news articles is to subscribe to various topic-specific newsfeeds and read what is automatically collated by their computers [Wang et al., 2006].

Although RSS—which is a subset of XML—is standardised², the practice of feed categorisation is not, meaning the granularity of topics which can be subscribed to is dependent on the publisher. This issue was addressed by Liu, Han, Noro and Tokuda [2009], with the design of a system which could essentially split or join existing RSS feeds to synthesise new ones based on user-specified keywords and queries.

Despite its shortcomings, RSS remains the most universal option for accessing feed content from a wide variety of news producers [O’Shea and Levene, 2011].

²<http://cyber.harvard.edu/rss/rss.html>

1.4.2 Keyword Extraction

Extracting relevant keywords from documents is not a new domain of research. Various methods have been presented, the most well-known being the intuitively logical tf-idf (term frequency, inverse document frequency) [Salton and Buckley, 1988] which ranks the significance of a term t in a document d which belongs to a corpus C as follows:

$$\text{tf-idf}(t) = \frac{\text{Occurrences}(t, d)}{\text{WordCount}(d)} \times \log_e \left(\frac{|C|}{|\{c \in C \mid t \in c\}|} \right) \quad (1.6)$$

tf-idf will extract the most unique keywords from a document within a corpus, because it penalises words which are common to many documents. However, in the context of a corpus of news articles, this uniqueness can lead to significant topic keywords being ignored because they appear with such frequency.

Bun and Ishizuka [2002] found that for news archive keyword extraction, a better alternative to tf-idf is tf-pdf (term frequency, proportional document frequency) as it is not biased against frequently repeated keywords.

Using tf-pdf, articles are modelled as belonging to one of a finite number of sources or *channels* within a corpus. The weighting of a term from an article within a channel is in this case linearly proportional to its frequency in the channel and exponentially proportional to the number documents in the channel where it occurs. A term's total weighting is the sum of its weightings across all channels, as can be seen in Equation 1.7 [Bun and Ishizuka, 2002], where:

- D = The number of channels in the corpus;
- K_c = The total number of terms in channel c ;
- F_{tc} = Frequency of term t in channel c ;
- n_{tc} = The number of articles in channel c where term t occurs;
- N_c = The total number of articles in channel c .

$$\text{tf-pdf}(t, D, K) = \sum_{c=1}^{c=D} \frac{F_{tc}}{\sqrt{\sum_{k=1}^{k=K_c} F_{kc}^2}} \times \exp \left(\frac{n_{tc}}{N_c} \right) \quad (1.7)$$

The reliance of both tf-idf and tf-pdf on a fixed background corpus results in a need to recompute the function for every document if any are added to or removed from the collection. This is impractical for large collections, and even in the case of large fixed collections it does not scale well, which has resulted in the development of other methods.

An approach derived from energy levels in quantum systems was proposed in Carpena et al. [2009], where keywords were extracted based on their spatial distributions within a single text. The theory behind the approach is that typically, keywords occurrences are distributed in significant frequency clusters throughout a document, whereas non-relevant words are distributed with uniform frequency (see Figure 1.7).

This technique allows relevant keywords to be distinguished from non-relevant common words with similar total frequencies without the use of a background corpus for comparison.




Figure 1.7: Frequency spectra for ‘Quixote’ and ‘but’ in the first 50,000 words of *Don Quixote*. [Carpena et al., 2009]

Several important observations have been made regarding keyword extraction for news articles specifically. Firstly, that important phrases in text are likely to be references to people, places and other named entities [Teitler et al., 2008]. Libraries such as the Stanford Named Entity Recognizer (NER) [Finkel and Manning, 2009] exist to extract these from text.

Secondly, that while 30% of an article’s keywords are inferred and cannot be found within the text without intelligent input, 60% are present in the article’s title and first few sentences [Lin and Hovy, 1997], since important facts are generally stated as part of an article’s *above the fold* content.

1.5 Evaluation Methods

Shahaf et al. [2012b] evaluated their system both for accuracy and with a user study, though as previously discussed they did not evaluate the aesthetic properties of the generated maps. The accuracy evaluation tested whether the system included the most ‘important’ (as decided by experts) documents in the map.

The user study focused on the strength of the results returned by specific queries, where output was transformed into a structureless list in order for the study to be double-blind against the other methods. The evaluation was performed between-subjects, so background knowledge had to be controlled for. Output was compared with that from Google News and a TDT (Topic Detection and Tracking) method presented in [Nallapati, 2003].

This approach to evaluation is less relevant to my proposed system, since it was actually evaluating the performance of the system in selecting documents based on a query, rather than visualising the documents on a map. In contrast, a visualisation and its usability is precisely the aspect of my process which I would need to evaluate.

The evaluation of TIARA [Liu, Zhou, Pan, Qian, Cai and Lian, 2009] was a more relevant method than the Metro Map evaluation as it was conducted against a baseline system which did not share any of its advanced features, although it was tailored for the same task; email

analysis. A series of questions were asked of participants, who used either TIARA or the baseline system to answer. The response time and accuracy of the participants was recorded, as well as their levels of satisfaction after completing the task. This evaluation used between-subjects designs and therefore required the use of a different dataset for each task, as the nature of the sensemaking means any repetition of the evaluation task on the same data would see participants' performance improve significantly due to recall alone.

1.6 Summary

To recap, this review began with an exploration of the problem of news overload using the findings of The Associated Press and Context-Based Research Group [2008], a field study conducted into the news consumption habits of young people. The findings of the study were then discussed, firstly in the context of four dimensions of information overload [Ho and Tang, 2001, Bergamaschi et al., 2010] and secondly in terms of how they relate to sensemaking; the cognitive process this project aims to support.

Information visualisation was identified as one common approach to both support sensemaking and reduce information overload, so various methods for visualising text-based documents and news articles specifically were presented and compared [Goyal et al., 2013, Havre et al., 2002, Liu, Zhou, Pan, Qian, Cai and Lian, 2009, Singh et al., 2015, Wang et al., 2006] as well as a more in-depth exploration of visual metaphors, in particular timeline, and schematic map visualisations. This section also featured the introduction of the metro map in its original context, before moving on to its use as a metaphor for visualisation in various domains.

Centrally, the work of Shahaf et al. [Shahaf and Guestrin, 2010, Shahaf et al., 2012*b,a*, 2013] which is particularly relevant to the aims and objectives of this project was discussed in detail, with an explanation of key metrics (*coherence*, *coverage* and *complexity*) defined in [Shahaf et al., 2012*b*] which are applicable to all graph-based representations of document collections. My main critiques of this body of work were the fixed background corpus and the physical layouts of the drawn maps when evaluated against the criteria defined in [Purchase et al., 1997, Stott et al., 2011].

Next followed a practical overview of methods for transforming news articles into entities with visual or comparative properties, including RSS feed mining, entity recognition, and keyword extraction. Lastly, approaches to experimental design and evaluation by several of the aforementioned studies were discussed in the context of their relevance to the proposed system.

The background, techniques and terminology discussed in this review will be taken forward into the next chapters, as they were all at least partially influential in the design and implementation of the system.

Chapter 2

Requirements

2.1 Project Scope

For the purpose of this dissertation and its timescale, the scope of the project was limited to the development of a system which generates metro maps based on the content of user-specified RSS feeds.

The subject of news fatigue and its possible remedies span multiple domains from data science to cognitive psychology to journalism and content publishing, so while there are many possible techniques in the feasible scope of the project, the practical scope required significant narrowing.

No focus was be given to the content of the RSS feeds themselves, as the standard is sufficiently specified in order for its attributes to be usable without understanding what is represented. Although it is a recognised problem that there is no standardisation for the granularity of RSS feed categories [Liu, Han, Noro and Tokuda, 2009], the only effect varying this would have on the system would be the granularity of the output visualisation.

Likewise, RSS feed discovery and recommendation, while clearly a potential extension to this work from the perspective of improving usability, was not considered.

In respect to the dimensions of information overload [Ho and Tang, 2001] discussed in the previous chapter; while a degree of attention will be given to addressing all four dimensions, the focus will be on information quality—in particular, contextual quality—and fragmentation. Information format, albeit an important facet of overload in general, does not vary significantly between major news publishers. Therefore, while this project may contribute a new unified format for displaying a collection of articles, changing the format or textual content of the articles themselves is not within its scope; any changes which do occur are incidental.

The aim of this work was ultimately the implementation of an end-to-end process for transforming news feeds into metro map schematics; that is to say, none of the techniques themselves are new. The area of interest was the transformation of data from articles to points in physical space on a transit map, and how varying the functions which compose the transformation affect the end result.

2.2 Requirements Gathering

This project was primarily research-based, so while I chose to follow typical software engineering practices by writing a requirements specification for organisational purposes, I did not undertake a formal requirements gathering process with potential users.

Instead, my requirements were derived from the work discussed in my literature review; features of the existing news visualisation system [Shahaf et al., 2012b], the transit map aesthetic principles formulated in [Stott, 2008, Stott et al., 2011], and Shneiderman’s [1996] infovis task taxonomy. Underpinning the technical requirements are the high-level recommendations for news producers specified by The Associated Press and Context-Based Research Group [2008] and Nordenson [2008]. The full requirements specification can be found in Appendices A.1 and A.2, with a discussion on certain conflicts and important decisions later in this chapter.

2.3 Prioritisation

Requirements were assigned priorities using the MoSCoW technique, since the size of the proposed system was not large enough to warrant more granularity in requirement priority. MoSCoW assigns requirements to one of four categories [Waters, 2009];

- **Must have** - Essential features required for the project to be useful.
- **Should have** - High value but non-critical features.
- **Could have** - Desirable features which will be moved out of scope if necessary.
- **Won’t have** - Features which have been requested but won’t be included.

As the requirements gathering process was based on analysing the findings of other researchers rather than surveying potential users, there were no requirements with a *won’t have* modifier.

2.3.1 Categorisation

The operation of the proposed system suggests a natural pipeline of four components through which data will be transformed, with each component comprising some distinct functionality which can be designed, implemented and if necessary modified, in isolation. The components are as follows;

- *Article Retrieval* - The process of parsing an RSS feed and downloading content from the articles it syndicates.
- *Keyword Extraction* - The language processing component, wherein articles are tokenised and their significant keywords are extracted.
- *Graph Building* - The transformation of a collection of articles and their associated keywords into a graph structure, by selecting keywords which best represent the entire corpus. The graph has no physical layout during this stage of the pipeline.
- *Map Drawing* - The generation of a visual representation of the graph structure in the form of a metro map, which the user will interact with.

In addition to the four stages of the pipeline, the system requires an ancillary storage component, to allow processed corpora and their graphs to be imported and exported. In the specification, all functional requirements were grouped according to one of the four stages, to assist in the implementation planning and testing processes.

2.4 High-Level Design

Figure 2.1 outlines the decomposition of the four pipeline components into their most important subtasks, with descriptions of the data transformations which occur at each stage.

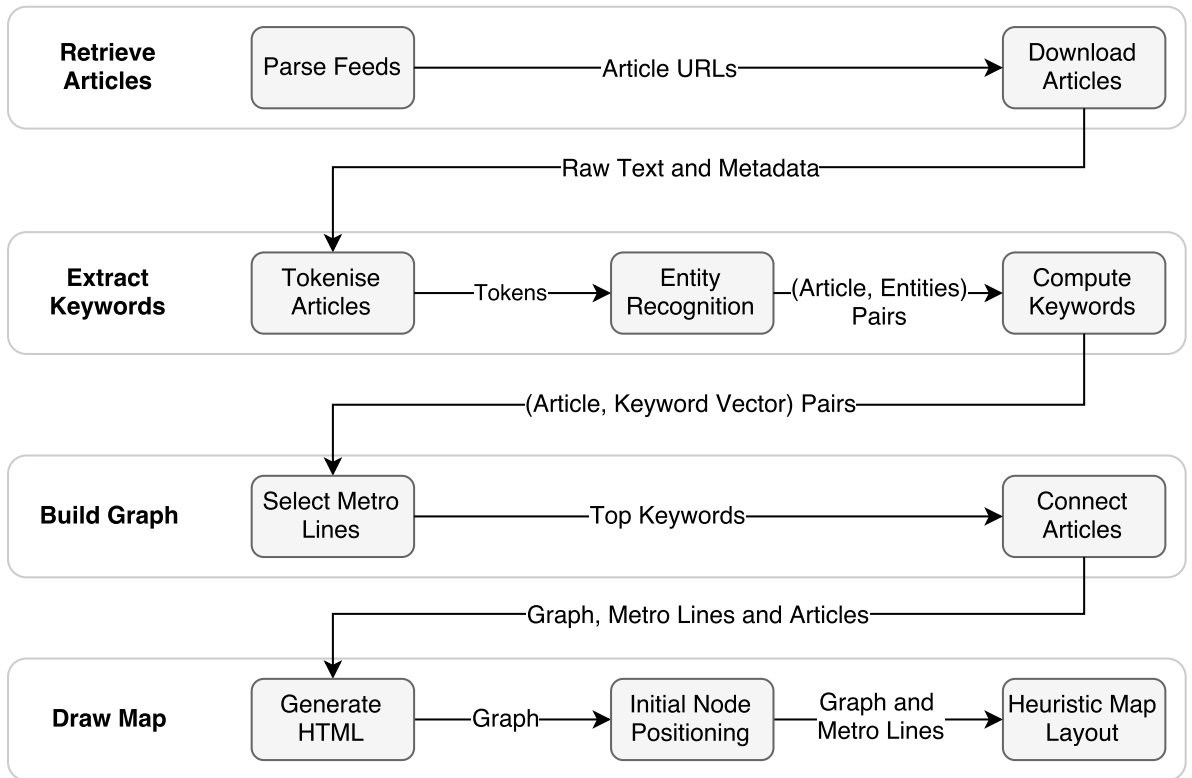


Figure 2.1: A conceptual model of data flow between components of the system.

Not included in Figure 2.1 is the boundary between the *Build Graph* and *Draw Map* stages of the pipeline, which marks the transition from the run-time Python which extracts and processes the article data, to JavaScript which only runs once the visualisation is opened in the browser; the entirety of the *Draw Map* stage. The boundary takes the form of an intermediate stage between *Build Graph* and *Draw Map*, where the topological map and the article data required for the visualisation are serialised to JSON¹.

¹JavaScript Object Notation; <http://www.json.org>

2.5 Discussion

Central to the importance of the project is the knowledge that news consumers do not subscribe to single RSS feeds; they specifically seek out software which aggregates multiple feeds for convenience [Wang et al., 2006]. Consequentially, F1.1 specifies that the system must accept multiple RSS feeds. However, news preferences are diverse, which gives rise to the potential case where a user specifies multiple feeds which do not share any significant keywords, and articles from one feed are excluded from the visualisation due to their lack of connectivity with the others.

Initially, I considered a requirement for the connectivity [Shahaf et al., 2012b] of every metro line to be greater than one; that is, no line should be included in the visualisation unless it intersects with another. However, in the case described above, mutually exclusive RSS feeds could lead to “orphaned” lines which still contain useful content but do not share contextual links to the main body of the map. The stories on these lines should not be ignored by the system, as they most likely would not be ignored by a reader using a traditional RSS reader. My argument is that the lack of connectivity of an article can itself be viewed as metadata on that article, and is not something to be penalised. However, a map exclusively containing orphaned lines is indicative of poor line selection, or—though less likely—a feed containing a set of completely unrelated articles. This highlights a fundamental assumption of the project; that in order to be of use, RSS feeds should contain explicitly related content.

A second issue of deliberation was the specification that stations on the generated maps should not be labelled (F4.7). The use of node labels is logical for a search task, where locating a particular entity or route on a map is the goal. However, when the goal is gaining an insight into the structure of the collection being visualised, labelling each station with the title of the article it represents would only reintroduce the information overload we have been trying to mitigate. The *overview* in the “Overview first” clause of the information seeking mantra [Shneiderman, 1996] is by definition the highest level of abstraction available on each data point, and while in a typical RSS reader this may well be the title, in our system this will not be the case. The context of the articles; i.e. the lines themselves, provide the overview, and the titles instead fall under “details-on-demand.”

One requirement which was left deliberately vague was (F4.8); where possible, maps should comply with Stott’s [2008] aesthetic criteria for metro map layouts. The lack of strictness at this stage was due to the non-spatial nature of the data being represented; while it is likely that there is an optimal layout for metro maps schematising real transit networks due to the geographic positioning of stations on those networks, there is no “sensible” underlying topology to our data. It was unknown at this stage whether we would be able to enforce strict layout criteria without significant pruning of the data.

Chapter 3

Design and Implementation

3.1 Introduction

During the requirements gathering process, four distinct components of the system were identified which form a pipeline of execution; *Article Retrieval*, *Keyword Extraction*, *Graph Building* and *Map Drawing*. Crucially, each is specified as being modular, allowing both for flexible extensibility and for alternative implementations to be tested directly against each other without requiring changes to the other components.

This chapter further decomposes these components to provide a detailed overview of the methodology used to implement them, followed by a discussion of the significant challenges and successes which arose during development.

3.2 Code Reuse

I developed the design of the system in Python 2.7 and JavaScript 1.7 using various open-source libraries and APIs for its ancillary functionality. The most notable are detailed below and discussed in context in the following sections.

- **FeedParser**¹: A Python module for downloading and parsing RSS feeds.
- **Newspaper**²: A Python library for downloading and extracting content and metadata from online articles.
- **lxml**³: A Python library for generating, parsing and manipulating XML and HTML.
- **NLTK (The Natural Language Toolkit)**⁴: A Python library for natural language processing and analysis.
- **Google Knowledge Graph Search**⁵: The API for Google's knowledge base, which returns structured semantic search results.
- **D3.js**⁶: A JavaScript library for creating and manipulating interactive web visualisations and their underlying data.

¹<http://pythonhosted.org/feedparser>

²<http://newspaper.readthedocs.io>

³<http://lxml.de>

⁴<http://www.nltk.org>

⁵<https://developers.google.com/knowledge-graph>

⁶<http://d3js.org>

3.3 Article Retrieval

The first stage of article retrieval is the parsing of RSS feeds, in order for article URLs and metadata to be extracted. The Python library FeedParser was used, as at the time of writing it provided the best support for RSS 0.9x, 1.0 and 2.0. The parsing process extracts a link, channel name, and the parsed publish date from every article in the feed, but it will also attempt to extract the author name if the `author` attribute is found.

Once the feed data has been extracted, it is used to construct an instance of `ArticleCollection`, which acts a wrapper around the contents of one or more feeds and provides the mechanism necessary to perform corpus-wide queries. The `Article` class encapsulates the functions for computing article-specific terms such as term-frequency for tf-idf, and the term-weighting component of tf-pdf. The result of this pre-processing stage is an `ArticleCollection` containing one or more serialisable `Articles` from one or more RSS feeds, where each `Article` contains extracted but unparsed data and associated metadata.

3.4 Keyword Extraction

The keyword extraction stage begins with the process of tokenising the body text of every article. Tokenisation here has three substages, all of which were implemented using NLTK (The Natural Language Toolkit for Python), and which are as follows:

1. Sentence segmentation: Split the text into a list of sentences and remove sentence punctuation.
2. Tokenisation: Split each sentence into a list of individual words and remove both whitespace and clause punctuation.
3. Part-of-Speech (POS) tagging: Categorise each token according to its lexical class, e.g. adjective. This more of an extension to the tokenisation process than a part of it, but it is performed directly after tokenisation and is necessary for the next stage of processing.

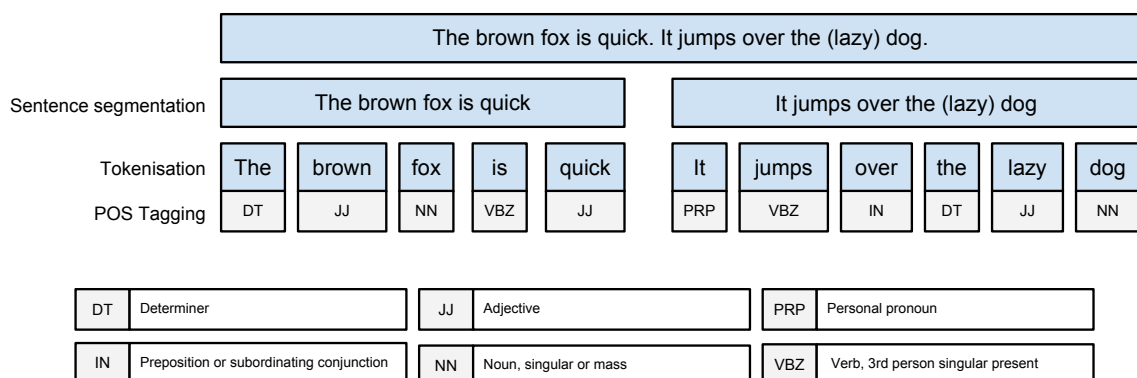


Figure 3.1: Stepping through the tokenisation process

3.4.1 Named Entity Recognition

The second half of the keyword extraction process is exclusively concerned with named entities within articles. This is because, intuitively, the names of people, places, events companies and other *things* form a set of strong candidate keywords. Restricting the candidates to entities alone reduces the search space by an order of magnitude when using frequency-based methods for keyword extraction, and bypasses the need for other natural language processing tasks such as stop-word removal and lemmatisation.

Once tokens have been POS tagged, named-entity chunking can be performed, again using NLTK. This process groups tokens into contiguous and non-overlapping *chunks*, where each chunk is a named entity; typically a proper noun or some other noun phrase. It is possible for chunks to contain other chunks (consider the chunk ‘Bank of England’, which contains the chunk ‘England’), but this is typically undesirable for entity recognition – where we desire specificity – so after chunking the tokens, we flatten the chunk structure so chunks cannot be any further decomposed.

At this stage, we have a list of chunks, each of which will be an eventual candidate for becoming a metro line. However, the chunks require some further processing before we can determine whether or not they are significant keywords.

Substring Matching

It is common stylistic practice in journalism to refer to the subjects of articles by their surnames. However, to avoid any confusion, the full names of those mentioned will often appear in the title or first few paragraphs of the article. If keyword strength is determined by frequency, then regardless of whether we use tf-idf or tf-pdf, we need every occurrence of an entity to refer to that entity by the same name; preferably the most specific, which is typically the longest.

During this stage of disambiguation, we only consider entities with more than one mention in the source text as candidates. This is because the chunking process can produce false positives by combining unrelated adjectives with valid chunks, but the likelihood of the same false positive being produced twice or more in the same article is low.

Let P denote the set of entities with more than one occurrence in the source article; S :

$$P = \{e \mid \text{occurrences}(e, S) > 1\}.$$

$\forall (e_1, e_2) \in \{P \times P \mid \text{length}(e_1) < \text{length}(e_2)\}$, if e_1 is a substring of e_2 , we call e_1 an *alias* of e_2 . Let A be the set of alias pairs in S ;

$$A = \{(a, e) \mid a \text{ is an alias of } e\}$$

We require the first term of every pair in A to be unique, but there is no such constraint for the second term (e); this allows multiple aliases to map to the same entity. For example,

$$\{('Zuckerberg', 'Mark Zuckerberg'), ('Zuck', 'Mark Zuckerberg')\}$$

would be unambiguous and therefore valid, but

$$\{('Mark', 'Mark Zuckerberg'), ('Mark', 'Zuckerberg')\}$$

would not be. Let U be the set of unambiguous alias-entity pairs in S :

$$U = \{(a, e) \mid (a, e) \in A \cap \forall (e_1, e_2) \in A, a = e_1 \implies e = e_2\}$$

We have a reasonable degree of confidence that for every alias-entity pair $(a, e) \in U$, occurrences of a in the source text can be replaced by e , making e a stronger candidate for keyword detection. Pairs in $(a, e) \in A \setminus U$, that is, aliases which could map to multiple entities in the source, are disregarded at this stage and left unchanged. Algorithm 1 illustrates how, given a list of entities, we can find the unambiguous pairs deterministically.

Algorithm 1: Finding unambiguous alias-entity pairs

Data: *names*: a list of recognised entities

Result: U : the set of unambiguous pairs in *names*

```

1  $U \leftarrow \{\}$ ;
2 foreach  $e_1, e_2 \in (\text{names} \times \text{names})$  do
3   if  $\text{len}(e_1) > \text{len}(e_2)$  then
4      $\text{swap}(e_1, e_2)$  ;
5   end
6   if  $e_1 \in e_2$  then                                     /* If  $e_1$  is a substring of  $e_2$  */
7     if  $e_1 \notin U$  then
8        $U[e_1] \leftarrow e_2$  ;                               /*  $(e_1, e_2)$  are a candidate pair */
9     else
10       $\text{delete } U[e_1]$  ;                                     /*  $e_1$  is now ambiguous, so remove it */
11    end
12  end
13 end

```

Entity Disambiguation with Knowledge Graph

The process described above is a simplistic approach which only addresses the matching of partial to full name mentions. In comparison, entity disambiguation describes the process of determining the identity of entities in a body of text. Performing this process on the sentence ‘The UK has voted to leave the EU,’ should identify ‘UK’ as ‘The United Kingdom’ and ‘EU’ as ‘The European Union’. As advanced methods for entity disambiguation are both complex and computationally expensive, I did not attempt to implement a formal method for this.

Instead, the list of recognised entities are queried against Google’s publicly accessible Knowledge Graph API, which returns a list of potential results as Schema.org⁷ types. Knowledge Graph is the service which replaced Freebase in 2015, and currently contains over 70 billion facts [Jeff Jarvis Twitter account, 2016].

⁷<http://schema.org> is an online hierarchy of types managed by W3C (The World Wide Web Consortium)

Dredze et al. [2010] identify three key challenges in entity linking using knowledge bases; name variations, ambiguity, and absence. Absence describes the lack of a corresponding entry for the entity in the knowledge base, which is not an easily tackled problem. Ambiguity is a consequence of the polysemy of many names and acronyms and requires disambiguation to be performed using more contextual methods.

Both of these problems are left unsolved in the design and implementation of the system due to scoping constraints. However, naive substring matching combined with the use of Knowledge Graph and some empirical parametric estimation are enough to disambiguate name variations of all forms (partial matches, abbreviations, and acronyms) to a sufficient degree, and with surprising accuracy.

The key to using Knowledge Graph for disambiguation lies in its most vaguely defined return value. Each result has a score attributed to it by Knowledge Graph, which is an indicator the strength of the match between the entity and the original query. Results are sorted by descending score; the higher the `resultScore`, the better the match. Although there is no defined upper limit for this value and no official documentation on how the score is derived, comparing the scores of the top two results for a query can provide a measure of certainty, for all but particularly esoteric or unknown entities.

We specify a threshold $\frac{1}{t}$, which is roughly proportional to the likelihood of accepting a false positive match. Then, if dividing the score of the first result by the score of the second yields a number greater than $\frac{1}{t}$, we accept the match. Provisionally we set $t = 0.5$; a discussion of how I arrived at this value is provided in the next section.

Using a knowledge base for disambiguation also inadvertently solves another problem I encountered while parsing articles, this time as a result the expository style of journalistic writing. While keywords are typically nouns or noun phrases, they can also appear in the form of denominal adjectives. These adjectives are derived from nouns; e.g. ‘French’ implies ‘France’ might be a keyword. Denominal adjectives are not amenable to traditional stemming or lemmatising, but querying Knowledge Graph for ‘French’ returns a top match of ‘France’ with a `resultScore` of 432.42807; more than four times larger than the next result.

Given a list of entities E , and top two results of a Knowledge Graph query for each $e \in E$; $R_e(1)$ and $R_e(2)$ with scores $S_e(1)$ and $S_e(2)$ respectively, our aim is to return a set of pairs mapping zero or more entities in E to their disambiguated forms;

$$K = \left\{ (e, R_e(1)) \mid \frac{S_e(1)}{S_e(2)} > \frac{1}{t} \right\}$$

Algorithm 2 describes this process. For higher values of t , the likelihood of accepting a false positive increases, and for lower values, the likelihood of accepting a false negative increases. For corpora which do not contain references to public figures and place names, choice of t should be empirically tuned against the prevalence of the expected entities.

It is unclear what should become of queries that Knowledge Graph only returns a single result for, but in this implementation they are ignored.

Algorithm 2: Entity disambiguation with Knowledge Graph

Data: *names*: a list of recognised entities
 t : the acceptance threshold for results, default = 0.5
Result: K : a set of disambiguation mappings for elements in *names*

```

1  $K \leftarrow \{\}$ ;
2  $T \leftarrow 1 \div t$ ;
3 foreach  $e \in names$  do
4    $results \leftarrow \text{KnowledgeGraphResults}(e)$ ;
5   if  $\text{len}(results) > 1$  then
6     if  $results[0].score \div results[1].score > T$  then
7        $K[e] = results[0].name$ ;
8     end
9   end
10 end

```

Determining Optimal Values for t

Empirically, I found the best values for t are in the range $0.35 < t < 0.65$, meaning we accept top results which are at least 1.5x-3x higher than the next best candidate. To determine this, I logged all the disambiguation pairs found in 20 articles from the BBC Politics RSS feed, letting t range over $\{0.3, 0.5, 0.7, 0.9\}$. I then manually classified the pairs as either true or false positives, where a true positive indicates a correctly identified entity, and a false positive is either an incorrectly identified entity or a non-entity which was not identified at all. Figure 3.2 shows the results of this investigation.




Figure 3.2: Disambiguation threshold against pairs found and false positives (%)

Although we see a reduction in false positives for smaller values of t , the trend-line illustrates that this is a game of diminishing returns; reducing t from 0.3 to 0.1 only reduces the false positives by 4.34%, but it leads to 43 fewer pairs being identified⁸.

There is a clear trade-off between the accuracy of the disambiguation and the number of false negatives we discard, but the cost of false positives in this case is less than the cost of false negatives. While a false positive could result in unrelated entities appearing as keywords for certain articles, the likelihood of the same false positive appearing with high frequency in enough articles to result in an erroneous metro line is incredibly low. In contrast, the cost of disregarding a false negative could result in articles being left off certain metro lines altogether. It is for this reason that we don't simply choose the value of t which yields the minimum ratio of false positives.

3.4.2 Choosing a Term-Weighting Metric for Keyword Ranking

Given a set of disambiguated entities for every article, the union of which forms a set of candidates keywords or *metro lines* for the collection, we must now determine which keywords are the most relevant. To do this, the two term-weighting methods described in Chapter 1 will be revisited; tf-idf [Sparck Jones, 1972] and tf-pdf [Bun and Ishizuka, 2002]. Both were implemented in the system so they could be compared directly against each another.

From an implementation perspective, the main difference between these two algorithms are the level at which they operate. tf-idf ranks keywords on a per-article basis, returning a vector of an article's keywords and their corresponding score against the background corpus. In contrast, tf-pdf is specified at corpus level and only returns a single metric for a given word; the sum of its significance across the whole corpus. In tf-pdf, corpuses have one or more *channels*, which contain articles. In our case, channels can be conveniently defined as RSS feeds from different publishers.

Bun and Ishizuka argue that tf-pdf is a more suitable metric for news corpora, because it discriminates between articles which originated from the same channel and articles in different channels, allowing keywords which are highly (and uniformly) frequent in one channel to be identified as significant in documents from a different channel. Similarly to tf-idf however, this process can be used to produce a vector of pairs of an article's highest ranked keywords. The actual scores attributed to keywords by either algorithm are not of direct importance, since it is only relative scores we use to construct the keyword vectors.

Regardless of which algorithm is used, the advantage to restricting the set of candidates to named entities becomes quickly apparent. Figure 3.3 shows the number of tokens against the percentage of extracted entities for 40 articles from the BBC's Politics RSS Feed.⁹ The interquartile range shows 50% of articles had entities comprising 5.4%-8.2% of their tokens after stop-word removal. With the mean equal to 6.65%, the implications of this are a search space which can be reduced by a factor of more than ten. Although performance optimisation was not specified in the aims of the project, gains of this nature are still significant.

At this point, given optional user-specified lists of keywords to either *include* or *ignore*, we

⁸Since it is both extensive and tangential to the focus of the project, raw data for this table can be found at <http://bit.ly/DisambiguationThresholding>.

⁹<http://feeds.bbc.co.uk/news/politics/rss.xml> (Accessed: 27/02/2017, Full data in Appendix B.1)




Figure 3.3: Entities as a percentage of Tokens across 40 BBC Politics articles.

can filter out or boost the scores of articles which mention these topics, which in the case of *include* will guarantee the keyword is selected as a metro line in the map, as long as it is mentioned in at least two articles.

The function of the *ignore* set is not to blacklist certain topics, as it will have no effect on articles which contain ignored keywords being placed on other metro lines. Instead, it performs domain specific stop-word removal for metro line candidates, preventing non-useful keywords with inevitably high scores (such as ‘United Kingdom’ for a metro map based on a UK news corpus) from being selected. During implementation, it was found to be useful to automatically append the names of every news publisher (e.g. ‘The Guardian’) in the feeds to the *ignore* set, as well as the names of the feeds themselves. This forces the system to select more specific candidates, even though their scores may be significantly lower.

While implementing the *ignore* functionality that the differing results of using tf-idf and tf-pdf on the same corpus became obvious. Because tf-idf penalises words which appear with low but constant frequency across a number of documents, keywords we intuitively would wish to ignore such as ‘United Kingdom’ are not scored as highly as they would be with tf-pdf. This is in line with the findings of Bun and Ishizuka; tf-pdf does do a better job at detecting low-frequency keywords which are consistent across many documents. However, this behaviour is exactly the opposite of what we desire from our metro lines; those low-frequency keywords are most likely already known or could be easily inferred by a user reading the map. Therefore, it is my recommendation that tf-idf be given preference in any implementations of this or similar systems, where the names of more specific low-level themes are being sought.

3.5 Graph Building

The key challenge in graph building is generalising the characteristics of a good overview, from both news consumer and metro map user perspectives. It is at this point that a significant difference between this system and the work of Shahaf et al. [2012a,b, 2013] becomes apparent.

Shahaf et al. generated maps in response to a query (such as the name of an event or time period, meaning the entire map would then be oriented around that query. In contrast, in our system the only query is an implied one; “What’s going on today?” which offers no starting point for choosing metro lines. Instead of a query-based search problem, we have a multi-document summarisation problem. While they were choosing a set of lines to fit a given query, we are concerned with choosing a set of lines which best cover an entire corpus.

3.5.1 Extending the Definition of Coverage to Paths

Our starting point is a corpus of articles and their associated keyword vectors, which form a metro map by Shahaf et al.’s definition, but one with too many edges and paths (metro lines) to be readable in practice. Therefore, in order to intelligently prune the number of paths to some fixed upper bound, we need a metric for comparing the relative importance of paths within a corpus.

To do this, we first recall the original definition of document coverage (equation 1.1 [Shahaf et al., 2012b,a, 2013]), where d is a document, w is a keyword, and \mathcal{W} is some normalised measure of term-frequency, such as tf-idf.

$$cover_d(w) : \mathcal{W} \rightarrow [0, 1] \quad (1.1)$$

In order to compare candidate metro lines, we build on this definition of coverage to extend it to paths (Equation 3.1).

$$Coverage(\mathcal{P}) = |\mathcal{P}| \sum_{d \in \mathcal{P}} \frac{cover_d(\mathcal{P})}{|\{p \in \Pi \mid d \in p, p \neq \mathcal{P}\}|} \quad (3.1)$$

Here, the extent to which a path \mathcal{P} covers a document d within a corpus \mathcal{D} is proportional to its tf-idf coverage of that document and inversely proportional to the number of other candidate paths which *also* cover that document. The coverage of the entire corpus by \mathcal{P} is then given the sum of its coverage of the articles along it multiplied by the length of the path, with high coverage being desirable.

It is important to note that although with the final multiplication we show preference to longer metro lines, the objective of the system is not simply to maximise the number of articles which are included on the generated maps. This could easily be achieved by choosing the most nonspecific universal keywords as metro lines, but the resulting map would be uninformative and would most likely do nothing to counter news overload. It is in the nature of summarisation that not all information can be preserved, and in doing this kind of graph selection, the information lost is those articles which don’t form links to the significant topics within the corpus.

3.5.2 Penalising Affinity

The principle of topic connectivity was the basis for choosing the metro map visualisation. Without it, we have simply generated a set of two-dimensional timelines with no contextual links. However, maps which are overly connected will quickly become unusable. In particular, maps where multiple lines runs adjacently through more than two nodes (see figure 3.4) a sign of poor line choice. This is not simply hyper-connectivity, but correlation between two or more lines, often resulting from keywords which are semantically close.




Figure 3.4: An example of three metro lines, each with high affinity.

We call this function of two lines their *affinity*, and define the affinity of a single line as the sum of its affinities over the set of paths, excluding itself.

$$Affinity(\mathcal{P}) = \frac{1}{|\mathcal{P}|} \sum_{p \in \Pi} 2^{|\{d \mid d \in \mathcal{P} \cap d \in p, p \neq \mathcal{P}\}|} \quad (3.2)$$

The affinity of two lines is defined as the number of articles they have in common, raised as a power of two. The effect of the power is illustrated in Figure 3.5. On the right, line *a* (orange) which shares one article with line *b* (green) and another with line *c* (blue) does not indicate as strong a correlation as on the left, where *a* shares two articles with *b* and none with *c*. As with line coverage, we penalise shorter lines.




Figure 3.5: ‘Run-on’ affinity (left) is penalised more heavily than ‘one-off’ affinity (right).

Since metro lines should have both high line coverage and low affinity, we can compute a score for every path by dividing the two composite measures (Equation 3.3).

$$LineScore(\mathcal{P}) = \frac{Coverage(\mathcal{P})}{Affinity(\mathcal{P})} \quad (3.3)$$

Sorting the list of candidates by descending score and filtering out those candidates which contain too few articles to be beneficial, we can then take the top n from the range $7 \leq n \leq 12$, to become the metro lines in our map. Any article which contains the name of a metro line in its keyword vector will be represented as a station, and articles with more than one will be represented as interchanges.

3.5.3 Coherence

The last metric proposed by Shahaf et al. is the one which presented the greatest conflict with the predefined scope of the system. For all but the most rapidly unfolding stories, forming a coherent chain of articles on one topic or event is not possible when only considering a single day's news. Shahaf et al. had the advantage of assuming any topic well-known enough to be queried against their system would most likely be made up of a coherent chain of events, such as the Greek debt crisis or Brexit. In contrast, at a daily summarisation level, often no such chains exist; many articles exist in isolation or instead represent the *beginning* of a potential future storyline.

Given a large enough background corpus, it would be possible to link the day's summary back in time to older articles, extending metro lines further back to support the *zoom and filter* or *details-on-demand* processes. A great deal of potential lies in this idea, as it could provide the story resolution the participants in the Associated Press and Context-Based Research Group [2008] study craved, allowing users to explore further back in time along individual metro lines while still centring itself around the most recent news. Unfortunately however, this is not something I could have achieved within the assigned timescale.

3.5.4 Serialisation

The intermediate stage of the pipeline between graph building and map drawing is purely an issue of implementation, with few design choices to be made. Starting with the Python representation of the metro map, we group articles by metro line to be serialised, along with their metadata and summaries, to JSON. This JSON is then copied into a preexisting JavaScript file which is responsible for representing and drawing the graph using D3.js, and the contents of the JavaScript file are embedded as a script in an HTML file, which contains general template markup for the page.

The main advantage to this process is that it results in a single file, meaning the resultant visualisations are as portable, shareable and platform independent as the news they were derived from. A secondary advantage is that the JSON accepted by D3 is standard and could easily be fed into a different JavaScript graphing library in the case of a more suitable alternative being written, with no change to the core Python software. The core software itself could also be hosted as a web app which serves the map HTML files, removing the need for client-side processing or installation.

3.6 Map Drawing

The metro map layout problem, that is, determining whether an optimal planar embedding exists for a map and a set of hard constraints which include drawing straight octilinear lines, was proven to be NP-Hard by Nöllenburg and Wolff [2005]. The authors present a mixed-integer program (MIP) approach to metro map layouts, which both guarantees an octilinear layout and avoids falling into local maxima (unlike Stott's), but not in polynomial time.

At the time of writing, there are no public domain algorithms or open-source libraries for generating automatic metro map layouts, in any language. Those which have been proposed in academia [Stott, 2008, Nöllenburg and Wolff, 2005] are too complex to implement given the current timescale, and all require NP-Hard optimisation techniques. To overcome this, a heuristic approach derived from the work of [Stott, 2008, Stott et al., 2011, Purchase, 1997] will be detailed.

In this section of the pipeline, we are no longer concerned with the contents of articles or the entities they reference. The map drawing context marks a shift in the terminology from the previous stages of the pipeline to the domain of graph drawing, but many terms have a one-to-one correspondence. A list of definitions is given below, and in addition we recall Shahaf et al.'s definition of a metro map from Chapter 1, which ties the terms together.

Definition 1. A metro map \mathcal{M} is a pair (G, Π) , where $G = (V, E)$ is a directed graph and Π is a set of paths, or metro lines in G . Each $e \in E$ must belong to at least one metro line.




Figure 3.6: Left-to-Right: A station (a), an interchange (b), a metro line (c), a link (d), a terminus (e) and a non-terminus station (f).

Station	A single article $v \in V$, represented by a node in the graph. If v has more than two incident edges, i.e an article which spans more than one topic in the graph, it is called an <i>interchange</i> .
Metro Line	A contiguous named path $p \in \Pi$ which is incident to two or more nodes, represented by a line of single colour. Metro lines are composed of <i>links</i> .
Link	A direct edge $e \in E$ between two nodes, which does not pass through any other node. In contrast to Shahaf et al.'s definition, here, links must belong to <i>exactly one</i> metro line. The case of two lines running side-by-side through the same two nodes is represented by two distinct links. This difference is due only to an implementation detail of the graphing library chosen.

Terminus A node with only one incident edge. This is not the definition of a terminus in the sense of the underlying metaphor, since by our definition, a station which is the terminus for two lines (see Figure 3.6.f) is no longer categorised as a terminus at all. Instead, the definition requires that the position of a terminus only determine the position of a single link in the graph.

We now turn our attention to transforming the metro map structure into a planar embedding of stations and the metro lines which connect them. As we have previously observed, not all graph layouts were created equal, and the aesthetic quality of a graph-based visualisations has a notable influence on their usability and ease of interpretation [Purchase et al., 1997, Purchase, 1997, Stott, 2008, Stott et al., 2011].

In Stott’s words: “[The] visualisation process is deeply rooted in the human interpretation of the graph as a structure and as such, the quality of the aesthetics of a particular drawing of a graph are very important.” [Stott, 2008, p.24]

We therefore devote a significant amount of thought to optimising the visual quality of the generated metro maps by revisiting Stott’s aesthetic criteria for layouts, discussed in Chapter 1. We justify which of these apply to our visualisations, and for those which are, we detail the methodology used to implement them.

3.6.1 Stott’s Criteria Metro Map layouts

The set of rules and criteria defined by Stott [2008] form a complex multicriteria optimisation function. Starting with an initial layout, nodes are embedded on a square grid of finite size, and iteratively moved to maximise the node movement criteria, as long as movement does not violate one of the node movement rules.

It should be noted that the function also applies to the positioning of labels and the maximum of a set of label movement criteria alongside the node movement criteria, but as visualisations are unlabelled, these will not be described.

As previously discussed, there are six node movement criteria, which – using the metro map-specific terms defined above – are as follows:

- **Angular Resolution:** Maximise the angle of incident links on the same metro line at each station. As Purchase [1997] found satisfying this criteria had no statistical effect on task performance, this will not be a priority¹⁰.
- **Edge Length:** Keep all links approximately the same length. This is a criterion which can be implemented by setting a soft constraint using the graphing library, but it will not be considered beyond this constraint.
- **Balanced Edge Length:** Keep links incident to a given station approximately the same length. While it would be possible to implement this as another soft constraint, there is no evidence that it would significantly improve usability, so this will not be considered.

¹⁰However, as angular resolution comes as a natural consequence of maximising line straightness, it is not disregarded completely during implementation.

- **Edge Crossings:** Minimise the number of link crossings. Purchase found this to be the most significant factor in determining task performance, therefore it will be considered.
- **Line Straightness:** Maximise the collinearity of links on the same metro line. Purchase found this criterion to be somewhat beneficial to task performance, so this is both a realistic and justified constraint which will be considered.
- **Octilinearity:** Draw links at multiples of 45° . Although Purchase found orthogonal graphs to be insignificant factors on task performance, there is a strong argument that the familiarity of the metro map metaphor hinges on certain similarities being preserved. Octilinearity is a noticeable design feature in the metro maps of the world, and therefore should be attempted in our visualisations.

We therefore consider minimising edge crossings to be the highest weighted criterion, followed jointly by maximising line straightness and octilinearity.

3.6.2 Initial Force-Directed Positioning

A major difference between the maps generated using this system and maps representing real transit networks is the lack of “real” starting positions we have for stations. While the freedom to move points around without having to preserve their relative positioning removes one layout constraint, finding a set of starting positions for the map layout which define a planar embedding with minimal line crossings is non-trivial.

As Stott notes, “This [difficulty] is due to our method being based on optimising an existing layout: if the initial embedding is not adequate, then our method may struggle to produce an acceptable optimisation.” [Stott, 2008, p.210]. One solution he suggests is to generate positions using a force-directed algorithm; a class of algorithms which simulate the forces of real physical motion to position nodes in a graph, with minimal edge crossings and approximately uniform edge lengths [Kobourov, 2012]. The downside to using a force-directed approach is that initial positioning and therefore the final map layout becomes nondeterministic.

D3.js includes an optimised implementation of Dwyer’s [2009] force-directed layout algorithm, and provides several parameters which can be used to alter the forces applied to some or all of the nodes in the graph. The two forces we will adjust to optimise the force-direction for metro maps are as follows:

- `force.charge([charge])`¹¹

Charge specifies the force of attraction between nodes, with negative values causing nodes to repel each other. In order to minimise crossings, this should be a large negative value, particularly for nodes which of higher degrees, which will be surrounded by more edges. We therefore set charge to -500 for nodes with weight > 1 , and -300 otherwise.

- `force.alpha([alpha])`¹²

Alpha is the graph’s cooling parameter, decaying as the simulation converges on a stable layout. As the temperature falls, node movement slows down, and eventually alpha will drop below a threshold which pauses the simulation, preventing further movement.

¹¹<https://github.com/d3/d3-3.x-api-reference/blob/master/Force-Layout.md#charge>

¹²<https://github.com/d3/d3-3.x-api-reference/blob/master/Force-Layout.md#alpha>

Unfortunately, the implementation of `d3.layout.force` lacks any further mechanism for detecting or reducing line crossings due to the computational expense of the best known algorithms (this subproblem is also NP-Hard). Even if line crossings could be detected, it is likely that the topology of certain graphs generated would make it impossible to find a planar embedding with no crossings.


The parameter `force.alpha` is re-evaluated every time the simulation is advanced by a single step, which happens approximately 60 times per second. Since our aim is to generate an initial layout with minimal line crossings which will result in the best possible map, we can constantly vary `force.alpha` based on the quality of the map as its stations are repositioned. Therefore, if the stations settle into a map layout with desirably high octilinearity and good line straightness, the energy of the graph will be reduced such that it is impossible for the embedding to change. Likewise, if the stations start repelling each other in a way that worsens line straightness and octilinearity, the system will have its energy increased and will be more likely to move out of the undesirable state.

The measure of quality that we use to adjust `force.alpha` is function of the two remaining chosen aesthetic criteria; octilinearity and line straightness. The means by which these two properties are calculated are described in the following sections.

Octilinearity

Octilinearity (Equation 3.4 [Stott, 2008]) penalises lines which are not drawn at multiples of 45° angles; the further the gradient of a link in the graph is from a multiple of 45° , the higher its octilinearity. Lower values are desirable, with zero indicating a graph is entirely octilinear.

$$\text{octilinearity}(G) = \sum_{u,v \in E} \left| \sin 4 \left(\tan^{-1} \frac{|u.y - v.y|}{|u.x - v.x|} \right) \right| \quad (3.4)$$



We can calculate the octilinearity of the graph in Figure 3.7 (left) as follows, taking positions relative to (0,0) in the bottom left corner. Algorithm 3 describes this calculation.

$$\text{octilinearity}(a, b) = \left| \sin 4 \left(\tan^{-1} \frac{|1 - 1|}{|4 - 1|} \right) \right| = 0$$

$$\text{octilinearity}(b, c) = \left| \sin 4 \left(\tan^{-1} \frac{|4 - 1|}{|5 - 4|} \right) \right| = 0.96$$

$$\text{octilinearity}(a, b, c) = 0 + 0.96 = 0.96$$

Figure 3.7: Octilinearity Calculation

There is a caveat to this formula which goes unmentioned by Stott, but which is necessary to emphasise as a detail of implementation. In the case where $u.x - v.x = 0$, we take octilinearity to be zero, as this value indicates a straight vertical line. Otherwise, attempting to calculate the inverse tangent of a fraction whose denominator is zero (this is mathematically valid, as the principle value of $\tan^{-1} \infty$ is $\frac{\pi}{2}$), which would result in undefined behaviour.

Algorithm 3: Calculating Map Octilinearity

Data: \mathcal{M} : a metro map ($G = (V, E), \Pi$)
Result: $oct \in \mathbb{Q}, 0 \leq oct \leq \frac{24}{25}|E|$: the octilinearity of \mathcal{M}

```


1  $oct \leftarrow 0$  ;
2 foreach  $e \in E$  do
3    $\delta_x \leftarrow |e.source.x - e.target.x|$  ;
4    $\delta_y \leftarrow |e.source.y - e.target.y|$  ;
5   if  $\delta_x = 0 \cup \delta_y = 0$  then
6      $slope \leftarrow 0$                                 /* This avoids a potential division by zero. */
7   else
8      $slope \leftarrow \frac{\delta_y}{\delta_x}$ 
9   end
10   $\theta \leftarrow |\sin(4 \times \tan^{-1}(slope))|$  ;
11   $oct \leftarrow oct + \theta$  ;
12 end

```

Line Straightness

The line straightness (Equation 3.5 [Stott, 2008]) criterion is defined as the sum of every ‘turn’ a metro line makes between its first and last stations, where a turn is measured as the angle between the new link direction and the continuation of the previous direction. The global line straightness is simply the sum of the straightness of every metro line in the map, and as with octilinearity, lower values are preferred.

$$\text{lineStraightness}(G) = \sum_{p \in \Pi} \sum_{u, v \in p} \theta(u, v) \quad (3.5)$$



The line straightness of the graph in Figure 3.8 (left) is therefore $\alpha_1 + \alpha_2$, where the values of α_1 and α_2 can be found geometrically from the direction of the links (Algorithm 4).

Figure 3.8: Line Straightness Calculation

Empirically, we found the best results came from weighting octilinearity and line straightness equally when using them to set `force.alpha`, but this could be refined through further investigation. We take the sum of the natural log of both values, to ensure that small improvements to already low values for either criterion are rewarded.

Algorithm 4: Calculating Map Line Straightness

Data: \mathcal{M} : a metro map ($G = (V, E), \Pi$)
Result: $ls \in \mathbb{Q}, ls > 0$: the line straightness of \mathcal{M}

```

1  $ls \leftarrow 0$  ;
2 foreach  $p \in \Pi$  do
    /* Greedily iterate over all 3-tuples of adjacent stations on  $p$  */
3   foreach  $(a, b, c) \in p$  do
4      $v_1 \leftarrow \frac{\vec{ab}}{|\vec{ab}|}$  ;  $v_2 \leftarrow \frac{\vec{bc}}{|\vec{bc}|}$  ; /* Normalise the two vectors to unit length */
      $\theta \leftarrow \cos^{-1}(v_1.x \times v_2.x + v_1.y \times v_2.y)$  ; /* Find the angle between links */
      $ls \leftarrow ls + \theta$  ;
5   end
6 end

```

3.6.3 Heuristic Refinement

Once all stations have been assigned their starting positions, we can start incrementally moving them to improve the quality of the map. The first step is to discretise the search space for station positions, by *snapping* stations to intersections on a large grid. We found the optimal size for this grid to be $\frac{n}{4} \times \frac{n}{4}$ for a map containing n stations.

Figure 3.9 shows the *snap-to-grid* process applied to a small example. We snap stations in order of descending weight (number of incident edges), as typically the more edges a node is connected to, the more difficult it is to place. In the example, this means station c (with weight three) is moved first. There is contention for the second move, as both stations a and b have weight two, so we make an arbitrary decision as to which we move first, and the two moves are shown as one step. The final station to move is station d , after which we can mark all stations as successfully placed.

Stations whose nearest intersections are already occupied are left unplaced. This implementation contrasts slightly with Stott's approach, where stations are snapped *in order* of Euclidian distance to their nearest intersection, and snapped to the next-nearest if the closest intersection is already occupied. We found this implementation performed poorly for maps with several highly connected nodes, which is a common due to news providers often publishing review articles which span many of the lines on the map.

To determine the final station positions, our approach now diverges from Stott's into a heuristic method which has been refined for metro maps comprising between 10 and 50 stations. It is vital that metro lines with high affinity have been removed before this stage, or the following process will result in a map where nodes with high weights are occluded.

Although by definition the *snap-to-grid* gives a map near-perfect octilinearity, it can significantly worsen line-straightness, as shown in Figure 3.10. It would be preferable to have the circled station repositioned as the transformation shows, even though this sacrifices the uniform edge length of the map. We therefore make several passes over each metro line, moving stations to the midpoint of the line between the previous and next neighbouring stations if that line is octilinear.




Figure 3.9: An example demonstrating the *snap-to-grid* process




Figure 3.10: An octilinear move which compromises the edge-length criterion

In addition, a final pass checks for stations which are still unplaced at this point and performs the same midpoint movement regardless of whether the line between the next and previous points is octilinear. This is simply to try and ensure the unplaced point obscures and crosses as few other links as possible in what is most likely a densely packed area of the map. An example of this process is shown in Figure 3.11.




Figure 3.11: A non-octilinear move of an unplaced station

After the final pass, we recompute octilinearity and line-straightness, and count how many points were left unplaced; i.e. drawn directly on top of another point. If any of the numbers are unsatisfyingly high, the process can be restarted and the randomness of the initial force-direction will result in a different initial embedding.

3.7 Summary

We have given a detailed description of the system’s pipeline and the novel transformation of news data from XML elements in an RSS feed to raw text, to parsed keywords, to recognised entities, to nodes in a directed graph, to stations on a formatted metro map.

Several unexpected difficulties arose during implementation, in particularly the tradeoffs involved with entity disambiguation in the context of what makes a “useful” metro line, and the process of deciding which lines to select between two or more with high affinity.

However, the most significant challenge of implementation was undoubtedly the process of drawing usable metro maps. The conflict between edge crossings, line-straightness and octilinearity was omnipresent, and the final algorithm still falls into both local and global minima as it is unable to detect the edge crossings. This section of the pipeline would benefit greatly from further work, as it is a significant area of active research in its own right. As the expected datasets for the map are typically small (under 50 stations), a non-polynomial time algorithm for refining station positions would be feasible in terms of running time, and would doubtless improve the quality of the maps generated.

The results obtained by using the system to generate metro maps for various news corpora are presented in the next chapter. We take forward the successes of entity disambiguation and denominal adjective stemming, the concepts of line coverage affinity, and the difficulties of generating initial station embeddings as points of discussion.

Chapter 4

Results and Discussion

This chapter presents the results of applying our method to a selection of RSS feeds spanning various news categories from well known news;

- **Politics** from *BBC News*, where we attempt to determine whether the map provides good coverage of the day's main stories (Section 4.1).
- **Business** from *The Telegraph*, where we discuss best and worst case map layouts including common pitfalls (Section 4.2);
- **Sport** from *The Guardian*, where we provide a warning on the risks of misinterpretation and false inferences (Section 4.3).

4.1 UK Politics (BBC News)

The main topic threads for March 16th 2017 in British politics were Scottish First Minister Nicola Sturgeon's call for a second independence referendum, the government's U-turn on National Insurance increases for the self-employed in the Spring Budget, and The Queen granting Article 50 royal assent, therefore enabling Brexit negotiations to begin.



Figure 4.1: The front page BBC Politics article; drawn as (a) in Figure 4.2.

To determine whether the map captured day’s important political issues, we turn back to the native representation of the news. The BBC selected a piece on the proposed Scottish referendum as the front page article for their Politics page (Figure 4.1); a position which is generally indicative of importance.

In contrast, RSS feeds have no notion of relative content importance, so the system has no way of knowing this article should be included. Nevertheless however, when the article is parsed, its keyword vector contains $\{\text{‘Scotland’}: 1.09861\}$ as the top entity. This term-frequency was strengthened by the transformation of every instance of the denominal adjective *Scottish* into the noun form, *Scotland*. The concentration of articles which focus heavily on Scotland (its coverage score of 182.644 was the third highest of the corpus) means it is selected as a metro line, resulting in the article being added to the map as we had hoped (see Figure 4.2).




Figure 4.2: A metro map covering the BBC UK Politics RSS feed (March 16th 2017)

As well as a series of articles following the progression of the proposed referendum, the system also selects the other two main threads as metro lines. *National insurance* is selected despite its relatively short length due to its low affinity with the other lines. The system interprets

this low affinity as the line covering a set of articles which no other line has, which boosts its likelihood of selection.

Despite being entirely subsumed by the *Labour Party* line, *Jeremy Corbyn* is still selected for the map because of the strength of its keyword score within the articles which contained it. Intuitively, we would hope that users realise lines which are entirely subsumed by other lines are important in their own right, while still being closely related to their subsuming entity.

4.2 Business (The Telegraph)

The system also performs well for international business news, where entities are typically the names of companies and organisations as opposed to people. NLTK Struggles with place names which have containment issues, for example, ‘City of London’ is split into ‘City’ and ‘London’, and ‘Bank of England’ into ‘Bank’ and ‘England’. ‘London’ and ‘England’ are then penalised by tf-idf for being frequent among much of the corpus, so neither are selected as metro lines. In spite of this, the map is still largely usable and exhibits good coverage.




Figure 4.3: A metro map covering The Telegraph’s Business RSS feed (March 17th 2017)

From an aesthetic perspective, a common pitfall of our approach is visible on the *Bank* line (Figure 4.3). During repeated passes along each octilinear line, stations are moved to the midpoint of their immediately adjacent neighbours, which results in noticeably variable edge lengths. A fix for this would be to iterate between pairs of stations which represent the endpoints of a run of stations all with weight two and distribute the accumulated stations at uniform intervals between the two points (Figure 4.4).



Figure 4.4: The redistribution of unbalanced edges along a line

A second pitfall is the inability of the algorithm to smooth the *European Union* line in Figure 4.5 between two points which are not octilinear to each other. The shape of the *China* line here is also frequently observed, because the initial force-directed algorithm causes lines to form circular arrangements, which are then snapped to co-ordinates which form squares.




Figure 4.5: A poor embedding of the map in Figure 4.3

4.3 Sport (The Guardian)

Sports journalism was popular among the participants of the Associated Press [2008] study, in part because of the clear cut categorisation of teams within sports, within countries where applicable. Story resolution could be always achieved on a match-by-match basis, or over one or more seasons of predefined length.

Our system was designed for use with current events articles rather than sports content, but the results obtained through doing so bring to light some interesting issues which were not evident in previous examples, and therefore warrant discussion.




Figure 4.6: A metro map covering the Guardian Sport RSS feed (March 16th 2017)

If someone who knew very little about rugby were to look at the metro map on sports and see *Six Nations*, they would most likely (and correctly) assume that it is the name of a tournament with six participating countries. Using only the connections drawn on the map however, they may also assume that all five countries whose metro lines intersect with Six Nations are participants. This is incorrect; New Zealand, which intersects twice, is not.

This is a good example of a metro map which exhibits desirable aesthetic properties and good coverage, but which is ambiguous and therefore uninformative at a thematic level. The names of the majority of the metro lines are countries, which create widespread entity ambiguity problems; does *England* imply the football team or the rugby team? The answer to this

question is unfortunately both, as the system is unable to disambiguate. The risk of misinterpretation by a user leading to confusion and contributing further to information overload is of course ever present, but particularly noticeable in this example.

The quality of the metro lines on this map confirmed our suspicion that the system is not practical for use with corpora which contain a large number of ambiguous entities. That is not to say, however, that metro maps are inappropriate for all news feeds related to sport. Generating a metro map based on The Guardian’s Football RSS feed (Figure 4.7) does not lead to the same team name ambiguities, and therefore results in an unequivocally distinct set of metro lines.




Figure 4.7: A metro map covering the Guardian Football RSS feed (March 17th 2017)

News coverage for individual sports and leagues, it transpires, is very amenable to being represented in the form of metro maps. This is mainly due to keyword extraction being less complex in this domain; the names of most teams and sports personalities are easily identifiable with a knowledge base. To only select *teams* as metro lines, we could simply filter out other entities at the Knowledge Graph stage, using the `Thing > Organization > SportsOrganization > SportsTeam` type from Schema.org. The same principle could be applied to countries or athletes, in order for users to adjust the schema represented by the metro lines at a *zoom-and-filter* level [Shneiderman, 1996].

Generalising this relationship between well defined schemas and a well defined set of metro lines to domains other than sports would suggest that in order to improve line selection for any map, it may be beneficial to restrict line choice to those which fit certain predefined types, such as `Thing > Person`, `Thing > Organization`, and `Thing > Place`.

4.4 Summary

This chapter presented three different perspectives from which the results of the system can be analysed; by content, by layout, and by context. Firstly, we verified that our maps provide sufficient coverage of topics which were deemed significant by the original news source, and discussed a concrete example of two topics being selected as metro lines despite their high affinity with each other.

Secondly, we took a step back from content and looked at the maps solely in terms of successful planar embeddings. We provided examples of a best and worst case layout for a particular current events metro map and identified common weaknesses of the map layout algorithm, in particular the inability of the layout algorithm to smooth non-octilinear edges once stations have been marked as placed.

Finally, we used the system to generate a metro map for a corpus it was not designed to be used on; international sports. This demonstrated the need for common entities within a corpus to be identifiable without ambiguity, unlike for instance the names of international sports teams between different sporting disciplines. We did find however that the system generates useful maps for sports news corpora when restricted to individual sports, where there is less obvious name ambiguity.

Section 4.3 ended with a more general note on the utility of a schema during the graph formation process, where metro lines are chosen from a pool of candidate keywords. Being able to select or reject lines based on a schematisation of expected entities within a corpus was not something we had considered during the implementation of the system, as it only came to light by using the system on data it was not designed for. Nevertheless, schema-based metro line selection is a novel concept which merits further investigation in future systems.

While accurate topic coverage and octilinear embeddings may be sufficient to evaluate the keyword extraction and map layout processes respectively, they do not provide any real insight into usability. That is to say, we may have designed a system which generates *good* metro maps both contextually and aesthetically, but the maps themselves have not been shown to address information overload. The next chapter therefore presents this evaluation through a series of experiments.

Chapter 5

Empirical Evaluation

Chapter 6

Conclusions

Bibliography

- Associated Press and Context-Based Research Group [2008], ‘A new model for news: Studying the deep structure of young adult news consumption’. Accessed: 26/10/2016.
URL: <http://manuscriptdepot.com/edition/documents-pdf/newmodel.pdf>
- Barzilay, R., McKeown, K. R. and Elhadad, M. [1999], Information fusion in the context of multi-document summarization, *in* ‘Proceedings of the 37th annual meeting of the Association for Computational Linguistics on Computational Linguistics’, Association for Computational Linguistics, pp. 550–557.
- Bergamaschi, S., Guerra, F. and Leiba, B. [2010], ‘Guest editors’ introduction: information overload’, *IEEE Internet Computing* **14**(6), 10–13.
- Binh Tran, G. [2013], Structured summarization for news events, *in* ‘Proceedings of the 22nd International Conference on World Wide Web’, WWW ’13 Companion, ACM, New York, NY, USA, pp. 343–348.
URL: <http://doi.acm.org/10.1145/2487788.2487940>
- British Broadcasting Corporation [2013], ‘Tube 150th anniversary: How the underground map evolved’. Accessed: 4/1/2017.
URL: <http://www.bbc.co.uk/news/uk-england-london-20943525>
- Bruns, A., Highfield, T. and Lind, R. A. [2012], ‘Blogs, twitter, and breaking news: The produsage of citizen journalism’, *Produsing theory in a digital world: The intersection of audiences and production in contemporary theory* **80**(2012), 15–32.
- Bun, K. K. and Ishizuka, M. [2002], ‘Topic extraction from news archive using tf-pdf algorithm’, *Proceedings of the 3rd International Conference on Web Information Systems Engineering*.
- Burkhard, R. A. [2004], Learning from architects: the difference between knowledge visualization and information visualization, *in* ‘Information Visualisation, 2004. IV 2004. Proceedings. Eighth International Conference on’, IEEE, pp. 519–524.
- Carpna, P., Bernaola-Galván, P., Hackenberg, M., Coronado, A. and Oliver, J. [2009], ‘Level statistics of words: Finding keywords in literary texts and symbolic sequences’, *Physical Review E* **79**(3), 035102.
- Dredze, M., McNamee, P., Rao, D., Gerber, A. and Finin, T. [2010], Entity disambiguation for knowledge base population, *in* ‘Proceedings of the 23rd International Conference on Computational Linguistics’, COLING ’10, Association for Computational Linguistics, Stroudsburg, PA, USA, pp. 277–285.
URL: <http://dl.acm.org/citation.cfm?id=1873781.1873813>

- Dwyer, T. [2009], Scalable, versatile and simple constrained graph layout, in ‘Computer Graphics Forum’, Vol. 28, Wiley Online Library, pp. 991–998.
- Eppler, M. J. [2004], Visuelle kommunikation der einatz von graphischen metaphern zur optimierung des wissenstransfers, in ‘Wissenskommunikation in Organisationen’, Springer, pp. 13–31.
- Eppler, M. J. [2006], ‘A comparison between concept maps, mind maps, conceptual diagrams, and visual metaphors as complementary tools for knowledge construction and sharing’, *Information visualization* 5(3), 202–210.
- Ferrari, D. and Mezzalana, L. [1970], ‘A computer-aided approach to integrated circuit layout design’, *Computer-Aided Design* 2(2), 19–23.
- Finkel, J. R. and Manning, C. D. [2009], Nested named entity recognition, in ‘Proceedings of the 2009 Conference on Empirical Methods in Natural Language Processing: Volume 1-Volume 1’, Association for Computational Linguistics, pp. 141–150.
- Fischer, G. and Stevens, C. [1991], Information access in complex, poorly structured information spaces, in ‘Proceedings of the SIGCHI Conference on Human Factors in Computing Systems’, CHI ’91, ACM, New York, NY, USA, pp. 63–70.
- Garg, A. and Tamassia, R. [1994], On the computational complexity of upward and rectilinear planarity testing, in ‘International Symposium on Graph Drawing’, Springer, pp. 286–297.
- @GoogleTrends Twitter account [2016], ‘Tweet: +250% spike in “what happens if we leave the EU” in the past hour’, <http://twitter.com/GoogleTrends/status/746137920940056578>. Accessed: 27/10/2016.
- Goyal, R., Malla, R., Bagchi, A., Mehta, S. and Ramanath, M. [2013], Esthete: A news browsing system to visualize the context and evolution of news stories, in ‘Proceedings of the 22Nd ACM International Conference on Information & Knowledge Management’, CIKM ’13, ACM, New York, NY, USA, pp. 2529–2532.
URL: <http://doi.acm.org/10.1145/2505515.2508208>
- Hargreaves, I., Thomas, J., Commission, I. T. and Commission, G. B. B. S. [2002], *New news, old news: an ITC and BSC research publication*, Independent Television Commission.
URL: <https://books.google.co.uk/books?id=VZ-hPAAACAAJ>
- Havre, S., Hetzler, E., Whitney, P. and Nowell, L. [2002], ‘Themeriver: Visualizing thematic changes in large document collections’, *IEEE transactions on visualization and computer graphics* 8(1), 9–20.
- Ho, J. and Tang, R. [2001], Towards an optimal resolution to information overload: An intermediary approach, in ‘Proceedings of the 2001 International ACM SIGGROUP Conference on Supporting Group Work’, GROUP ’01, ACM, New York, NY, USA, pp. 91–96.
URL: <http://doi.acm.org/10.1145/500286.500302>
- Hochmair, H. [2009], ‘The influence of map design on route choice from public transportation maps in urban areas’, *The Cartographic Journal* 46(3), 242–256.

- Holton, A. E. and Chyi, H. I. [2012], ‘News and the overloaded consumer: Factors influencing information overload among news consumers’, *Cyberpsychology, Behavior, and Social Networking* **15**(11), 619–624.
- Husin, H. S., Thom, J. A. and Zhang, X. [2014], Analysing user access to an online newspaper, in ‘Proceedings of the 2014 Australasian Document Computing Symposium’, ADCS ’14, ACM, New York, NY, USA, pp. 77:77–77:80.
- @jeffjarvis Twitter account [2016], ‘Tweet: *Google knowledge graph has more than 70 billion facts about people, places, things. + language, image, voice translation*’, <https://twitter.com/jeffjarvis/status/783338071316135936>. Accessed: 20/02/2017.
- Kobourov, S. G. [2012], ‘Spring embedders and force directed graph drawing algorithms’, *CoRR* **abs/1201.3011**.
URL: <http://arxiv.org/abs/1201.3011>
- Lin, C.-Y. and Hovy, E. [1997], Identifying topics by position, in ‘Proceedings of the fifth conference on Applied natural language processing’, Association for Computational Linguistics, pp. 283–290.
- Lipton, R. J., North, S. C. and Sandberg, J. S. [1985], A method for drawing graphs, in ‘Proceedings of the first annual symposium on Computational geometry’, ACM, pp. 153–160.
- Liu, B., Han, H., Noro, T. and Tokuda, T. [2009], Personal news rss feeds generation using existing news feeds, in ‘International Conference on Web Engineering’, Springer, pp. 419–433.
- Liu, S., Zhou, M. X., Pan, S., Qian, W., Cai, W. and Lian, X. [2009], Interactive, topic-based visual text summarization and analysis, in ‘Proceedings of the 18th ACM Conference on Information and Knowledge Management’, CIKM ’09, ACM, New York, NY, USA, pp. 543–552.
- Nallapati, R. [2003], Semantic language models for topic detection and tracking, in ‘Proceedings of the 2003 Conference of the North American Chapter of the Association for Computational Linguistics on Human Language Technology: Proceedings of the HLT-NAACL 2003 student research workshop-Volume 3’, Association for Computational Linguistics, pp. 1–6.
- Nguyen, P. H., Xu, K., Walker, R. and Wong, B. [2014a], ‘Timesets: timeline visualization for sensemaking’.
- Nguyen, P. H., Xu, K., Walker, R. and Wong, B. W. [2014b], Schemaline: timeline visualization for sensemaking, in ‘Information Visualisation (IV), 2014 18th International Conference on’, IEEE, pp. 225–233.
- Nöllenburg, M. and Wolff, A. [2005], Automated drawing of metro maps, PhD thesis, Technical University of Munich. Accessed: 15/12/2016.
URL: <https://www.semanticscholar.org/paper/Automated-Drawing-of-Metro-Maps-Nllenburg-Wolff>
- Nordenson, B. [2008], ‘Overload! Journalism’s battle for relevance in an age of too much information’, *Columbia Journalism Review* **47**(4), 30.

- O'donnell, A. M., Dansereau, D. F. and Hall, R. H. [2002], 'Knowledge maps as scaffolds for cognitive processing', *Educational psychology review* **14**(1), 71–86.
- Old, L. J. [2002], Information cartography: Using gis for visualizing nonspatial data, in 'Proceedings of the ESRI International Users Conference', Citeseer.
- O'Shea, M. and Levene, M. [2011], 'Mining and visualising information from rss feeds: a case study', *International Journal of Web Information Systems* **7**(2), 105–129.
- Pentina, I. and Tarafdar, M. [2014], 'From "information" to "knowing": Exploring the role of social media in contemporary news consumption', *Computers in Human Behavior* **35**, 211–223.
- Pera, M. S. and Ng, Y.-K. [2008], 'Utilizing phrase-similarity measures for detecting and clustering informative rss news articles', *Integrated Computer-Aided Engineering* **15**(4), 331–350.
- Phuvipadawat, S. and Murata, T. [2010], Breaking news detection and tracking in twitter, in 'Web Intelligence and Intelligent Agent Technology (WI-IAT), 2010 IEEE/WIC/ACM International Conference on', Vol. 3, IEEE, pp. 120–123.
- Purcell, K., Rainie, L., Mitchell, A., Rosenstiel, T. and Olmstead, K. [2010], 'Understanding the participatory news consumer', *Pew Internet and American Life Project* **1**, 19–21.
- Purchase, H. [1997], Which aesthetic has the greatest effect on human understanding?, in 'International Symposium on Graph Drawing', Springer, pp. 248–261.
- Purchase, H. C., Cohen, R. F. and James, M. I. [1997], 'An experimental study of the basis for graph drawing algorithms', *J. Exp. Algorithmics* **2**.
URL: <http://doi.acm.org/10.1145/264216.264222>
- Rennison, E. [1994], Galaxy of news: An approach to visualizing and understanding expansive news landscapes, in 'Proceedings of the 7th Annual ACM Symposium on User Interface Software and Technology', UIST '94, ACM, New York, NY, USA, pp. 3–12.
URL: <http://doi.acm.org/10.1145/192426.192429>
- Rosen, J. [2008], 'National explainer: A job for journalists on the demand side of news', http://archive.pressthink.org/2008/08/13/national_explain.html. Accessed: 2/11/2016.
- Russell, D. M., Slaney, M., Qu, Y. and Houston, M. [2006], Being literate with large document collections: Observational studies and cost structure tradeoffs, in 'Proceedings of the 39th Annual Hawaii International Conference on System Sciences (HICSS'06)', Vol. 3, IEEE, pp. 55–55.
- Salton, G. and Buckley, C. [1988], 'Term-weighting approaches in automatic text retrieval', *Information processing & management* **24**(5), 513–523.
- Schick, A. G., Gordon, L. A. and Haka, S. [1990], 'Information overload: A temporal approach', *Accounting, Organizations and Society* **15**(3), 199–220.

- Shahaf, D. and Guestrin, C. [2010], Connecting the dots between news articles, *in* ‘Proceedings of the 16th ACM SIGKDD international conference on Knowledge discovery and data mining’, ACM, pp. 623–632.
- Shahaf, D., Guestrin, C. and Horvitz, E. [2012a], Metro maps of science, *in* ‘Proceedings of the 18th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining’, KDD ’12, ACM, New York, NY, USA, pp. 1122–1130.
URL: <http://doi.acm.org/10.1145/2339530.2339706>
- Shahaf, D., Guestrin, C. and Horvitz, E. [2012b], Trains of thought: Generating information maps, *in* ‘Proceedings of the 21st International Conference on World Wide Web’, WWW ’12, ACM, New York, NY, USA, pp. 899–908.
URL: <http://doi.acm.org/10.1145/2187836.2187957>
- Shahaf, D., Yang, J., Suen, C., Jacobs, J., Wang, H. and Leskovec, J. [2013], Information cartography: creating zoomable, large-scale maps of information, *in* ‘Proceedings of the 19th ACM SIGKDD international conference on Knowledge discovery and data mining’, ACM, pp. 1097–1105.
- Shneiderman, B. [1996], The eyes have it: A task by data type taxonomy for information visualizations, *in* ‘Visual Languages, 1996. Proceedings., IEEE Symposium on’, IEEE, pp. 336–343.
- Singh, J., Anand, A., Setty, V. and Anand, A. [2015], Exploring long running news stories using wikipedia, *in* ‘Proceedings of the ACM Web Science Conference’, WebSci ’15, ACM, New York, NY, USA, pp. 57:1–57:2.
URL: <http://doi.acm.org/10.1145/2786451.2786489>
- Sparck Jones, K. [1972], ‘A statistical interpretation of term specificity and its application in retrieval’, *Journal of documentation* **28**(1), 11–21.
- Stott, J. [2008], Automatic layout of metro maps using multicriteria optimisation, PhD thesis, University of Kent. Accessed: 11/11/2016.
URL: <http://www.jstott.me.uk/thesis/thesis-final.pdf>
- Stott, J., Rodgers, P., Martinez-Ovando, J. C. and Walker, S. G. [2011], ‘Automatic metro map layout using multicriteria optimization’, *IEEE Transactions on Visualization and Computer Graphics* **17**(1), 101–114.
- Strong, D. M., Lee, Y. W. and Wang, R. Y. [1997], ‘Data quality in context’, *Communications of the ACM* **40**(5), 103–110.
- Tamassia, R. [1987], ‘On embedding a graph in the grid with the minimum number of bends’, *SIAM Journal on Computing* **16**(3), 421–444.
- Teitler, B. E., Lieberman, M. D., Panozzo, D., Sankaranarayanan, J., Samet, H. and Sperling, J. [2008], Newsstand: a new view on news, *in* ‘Proceedings of the 16th ACM SIGSPATIAL international conference on Advances in geographic information systems’, ACM, p. 18.
- Transport For London [2014], ‘Harry Beck’s Tube map’, <https://tfl.gov.uk/corporate/about-tfl/culture-and-heritage/art-and-design/harry-becks-tube-map>. Accessed: 18/1/2017.

- Wang, T., Yu, N., Li, Z. and Li, M. [2006], nReader: Reading News Quickly, Deeply and Vividly, *in* ‘CHI ’06 Extended Abstracts on Human Factors in Computing Systems’, CHI EA ’06, ACM, New York, NY, USA, pp. 1385–1390.
- Waters, K. [2009], ‘Prioritization using moscow’, *Agile Planning* **12**.
- Weick, K. E., Sutcliffe, K. M. and Obstfeld, D. [2005], ‘Organizing and the process of sense-making’, *Organization science* **16**(4), 409–421.
- Wise, J. A., Thomas, J. J., Pennock, K., Lantrip, D., Pottier, M., Schur, A. and Crow, V. [1995], Visualizing the non-visual: Spatial analysis and interaction with information from text documents, *in* ‘Proceedings of the 1995 IEEE Symposium on Information Visualization’, INFOVIS ’95, IEEE Computer Society, Washington, DC, USA, pp. 51–.
- Yen, J.-C., Lee, C.-Y., Chen, I. et al. [2012], ‘The effects of image-based concept mapping on the learning outcomes and cognitive processes of mobile learners’, *British Journal of Educational Technology* **43**(2), 307–320.
- Yi, J. S., Kang, Y.-a., Stasko, J. T. and Jacko, J. A. [2008], Understanding and characterizing insights: How do people gain insights using information visualization?, *in* ‘Proceedings of the 2008 Workshop on BEyond Time and Errors: Novel evaluation Methods for Information Visualization’, BELIV ’08, ACM, New York, NY, USA, pp. 4:1–4:6.
URL: <http://doi.acm.org/10.1145/1377966.1377971>
- Ziemkiewicz, C. and Kosara, R. [2009], Preconceptions and individual differences in understanding visual metaphors, *in* ‘Computer Graphics Forum’, Vol. 28, Wiley Online Library, pp. 911–918.

Appendix A

Full Requirements Specification

A.1 Functional Requirements

1 Article Retrieval

Description

The article retrieval component of the system will accept the URL of one or more RSS feeds, collect links to the articles referenced and parse the text of those articles for further processing.

Functional Requirements

- F1.1 The system must accept any standard XML document compliant with the RSS 2.0 specification¹, i.e. it should not be specific to any particular news provider
- F1.2 The system must be able to extract a specified number of articles from an RSS feed, in reverse chronological order.
- F1.3 The system must be able to download the textual content and/or raw HTML for each article.
- F1.4 The system should accept multiple RSS feeds from one or more news provider(s) and merge their content into one collection.
- F1.5 The system could verify new article URLs against the URLs of imported articles to ensure no articles are duplicated.

2 Keyword Extraction

Description

The keyword extraction stage will tokenise the parsed article text and determine a set of significant keywords for each article individually.

¹<http://cyber.harvard.edu/rss/rss.html>

Functional Requirements

- F2.1 The system must tokenise articles in order to perform basic natural language processing such as stop-word extraction and lemmatising.
- F2.2 The system must implement a method for keyword extraction, and calculate a corresponding measure of relative importance for each keyword such as tf-idf.
- F2.3 The system could attempt to combine keywords it considers equivalent (e.g. *UK* and *United Kingdom*) to form stronger keyword matches between or within articles.
- F2.4 The system could use external services (e.g. Google’s Knowledge Graph API²) to query any extracted keywords, in order to gain further insight or perform entity disambiguation [Dredze et al., 2010].

3 Graph Building

Description

This process involves determining a set of corpus keywords from the union of all the articles’ keywords to form connected paths of edges (*lines*), and fitting a maximal number of articles into the resulting graph.

Functional Requirements

- F3.1 The system must analyse the keywords extracted from all articles in a corpus to choose a set of the n most significant topics, where n is either predetermined or user-specified.
- F3.2 The system must use the extracted topics and the publish dates (which form a natural ordering of nodes) of the articles to form a directed graph, with articles as vertices and common topic storylines as edges.
- F3.3 The system should accept a user-specified topic or list of topics to include or exclude from the graph.
- F3.4 The system should choose topics which are specific to some but not all articles in the collection, so as to avoid highly correlated topic keywords.
- F3.5 The system could support exporting generated graphs in a graph description language e.g. DOT³ or GraphML⁴.
- F3.6 The system could attempt to combine keywords to form topics if it considers them highly correlated.
- F3.7 The system could attempt to maximise the coverage of the topic selection, i.e. maximise the number of articles covered by a given set of keywords.

²<http://www.google.com/intl/es419/insidesearch/features/search/knowledge.html>

³<http://www.graphviz.org/content/dot-language>

⁴<http://graphml.graphdrawing.org>

4 Map Drawing

Description

The visualisation component will generate an interactive visualisation of the graph which can be used to explore the corpus as a whole and drill-down to the individual article level.

Functional Requirements

- F4.1 The system must provide the capability for users to visualise the graph structures it generates using any HTML5 compliant web browser.
- F4.2 Metro lines must each be drawn in a different colour which contrasts that of other lines.
- F4.3 The visualisation must include a key mapping metro line names to their colours on the map.
- F4.4 Stations must be drawn with common (non-unique) symbols, with a distinction made for interchange stations.
- F4.5 The system must provide drill-down details for nodes, e.g. by providing a hyperlink to the original article or embedding static content from each article within the visualisation itself to provide a preview.
- F4.6 The maps generated should be readable by ensuring nodes and edges do not overlap with each other where possible.
- F4.7 In support of F4.6, to improve the readability of the map, nodes should not be labelled with article titles.
- F4.8 The maps generated should adhere (where possible) to the metro map aesthetics defined in [Stott, 2008, Stott et al., 2011] to preserve the familiarity of the metaphor.
- F4.9 The system could allow some degree of interactive customisation which does not change the underlying structure of the graph, such as dragging nodes or changing attributes including colour.

5 Storage and Persistence

Description

This component of the system is responsible for saving and importing previously downloaded corpora and reconstructing their graphs.

Functional Requirements

- F5.1 The system must support the importing/exporting of graph and article data in an intermediate data form, in order to fully reconstruct graphs it had previously created.

F5.2 By default, the articles collected by each run of the system must be treated as a new corpus so keyword ranking is deterministic for any given feed.

A.2 Nonfunctional Requirements

1 Security

The system will not require any kind of authentication to use, and will only stores data which is publicly available. As there are no security regulations which govern its usage, security is not a critical consideration and there are only two associated requirements.

NF1.1 The system will not collect any data during installation and usage without obtaining consent from the user.

NF1.2 The system will not transmit any data which was necessary to collect or generate, including log files, without obtaining explicit consent from the user.

2 Software Quality

The following list specifies the system's core requirements in terms of portability, source control, testability, usability and documentation. There are no specific performance metric requirements for the system at this stage of its development.

NF2.1 The system must not use any platform-specific libraries, functions or commands.

NF2.2 The system must provide a `requirements.txt`⁵ file or similar, to allow its dependencies to be installed using Pip.

NF2.3 The system must be versioned and privately hosted on GitHub.

NF2.4 The implementation of the system should include a severity-based logging facility which writes to a text file, for use during debugging and testing.

NF2.5 The system must provide a non-interactive help facility for users.

NF2.6 The system should provide visual feedback during computationally expensive tasks to show task progress, e.g. with loading bars.

⁵<https://pip.readthedocs.io/en/1.1/requirements.html>

Appendix B

Implementation Raw Data

B.1 Token/Entity Data

Table B.1: Entities as a percentage of total tokens from 40 BBC Politics articles.

Tokens	Entities	Entities (%)
81	7	8.64198
129	8	6.20155
167	9	5.38922
181	15	8.28729
252	17	6.74603
264	15	5.68182
272	11	4.04412
275	17	6.18182
291	21	7.21649
343	22	6.41399
349	14	4.01146
366	23	6.28415
367	28	7.62943
367	14	3.81471
404	42	10.39604
409	40	9.77995
413	41	9.92736
434	25	5.76037
444	20	4.5045
453	29	6.40177
480	28	5.83333
498	41	8.23293
503	30	5.96421
523	35	6.69216
527	17	3.22581
536	62	11.56716
562	39	6.9395
562	32	5.69395

Table B.1: Entities as a percentage of total tokens from 40 BBC Politics articles.

Tokens	Entities	Entities (%)
571	42	7.35552
618	33	5.33981
621	38	6.11916
674	35	5.19288
688	34	4.94186
710	63	8.87324
722	65	9.00277
831	76	9.14561
920	51	5.54348
987	83	8.40932
1119	52	4.64701
1243	52	4.18343